

Final review

MICROSERVICES



Final


■ Dalby Hall

■ 9:30 – 12:00


- Closed book/notes.
- No personal items including electronic devices (cell phones, computers, calculators, PDAs).
- No additional papers are allowed. Sufficient blank paper is included in the exam packet.
- Bring only pen, pencil (eraser)




Architecture styles




Layering



Components




Service Oriented
Architecture



Microservice architecture

Event
Driven
Architecture




Stream based architecture




Service Oriented Architecture

- Disadvantages
 - Complex ESB
 - Changing the business process while still business processes are running is very difficult
 - Most SOA's are build on top of monoliths



Problem with SOA

- Most SOA's are build on top of monoliths




Problems with a monolith architecture

- Can evolve in a big ball of mud
- Limited re-use is realized across monolithic applications
- All or nothing scaling
- Single development stack
- Does not support small agile scrum teams
- Deploying a monolith takes a lot of ceremony



Microservices

- Small independent services
 - Simple and lightweight
 - Runs in an independent process
 - Language agnostic
 - Decoupled



Appropriate boundaries

- DDD bounded context
- Autonomous functions
- Size of deployable unit
- Most appropriate function or subdomain
- Polyglot architecture
- Selective scaling
- Small agile teams
- Single responsibility
- Replicability or changeability
- Coupling and cohesion



Orchestration vs. choreography

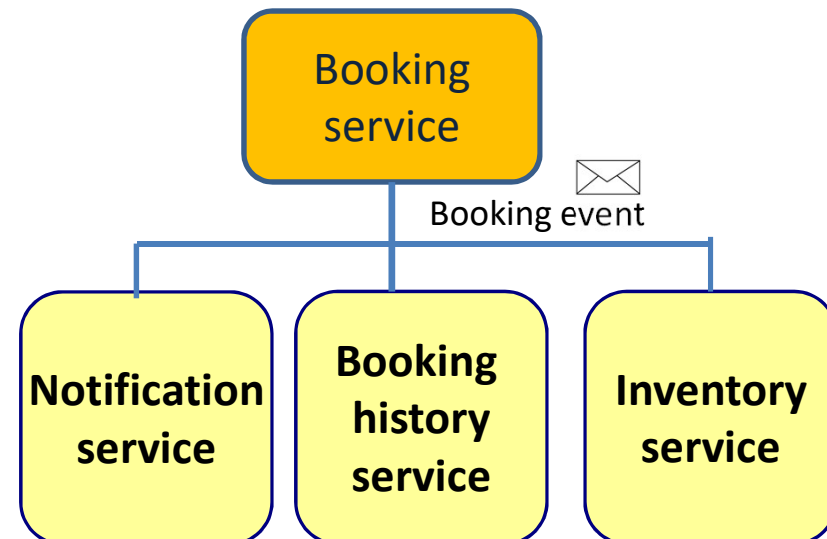

- Orchestration

- One central brain

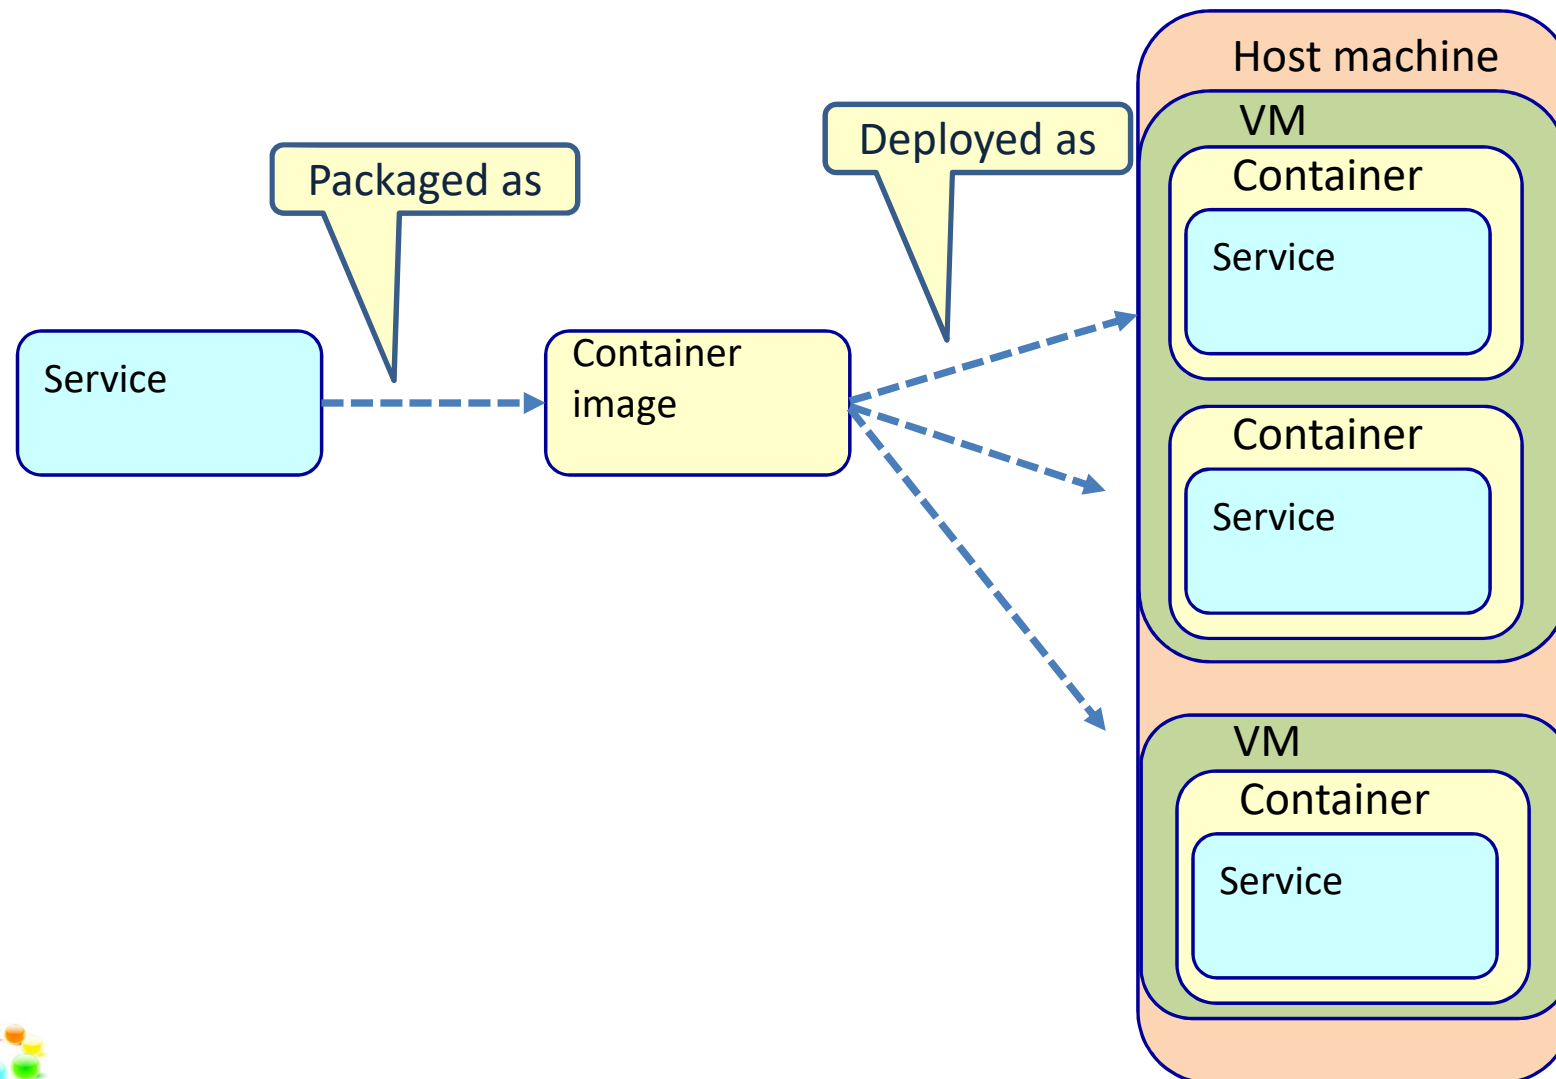


- Choreography

- No central brain



Service per container



From monolith to microservice

- Not a big bang
- Strangler approach
- Separate front-end and back-end
- Extract a service



Why microservices?

- Agility
 - Much easier to respond to change
- Testability
 - Easier to test
 - Scope is smaller
- Deployability
 - Less ceremony
 - Less risk
- Scalability
- Availability
 - Fault tolerance




CONFIG SERVICE




Spring cloud config

- File based configuration



- Git based configuration





NETFLIX EUREKA

SERVICE REGISTRY: EUREKA



Using Eureka




NETFLIX

RIBBON

LOAD BALANCING: RIBBON



Load balancer

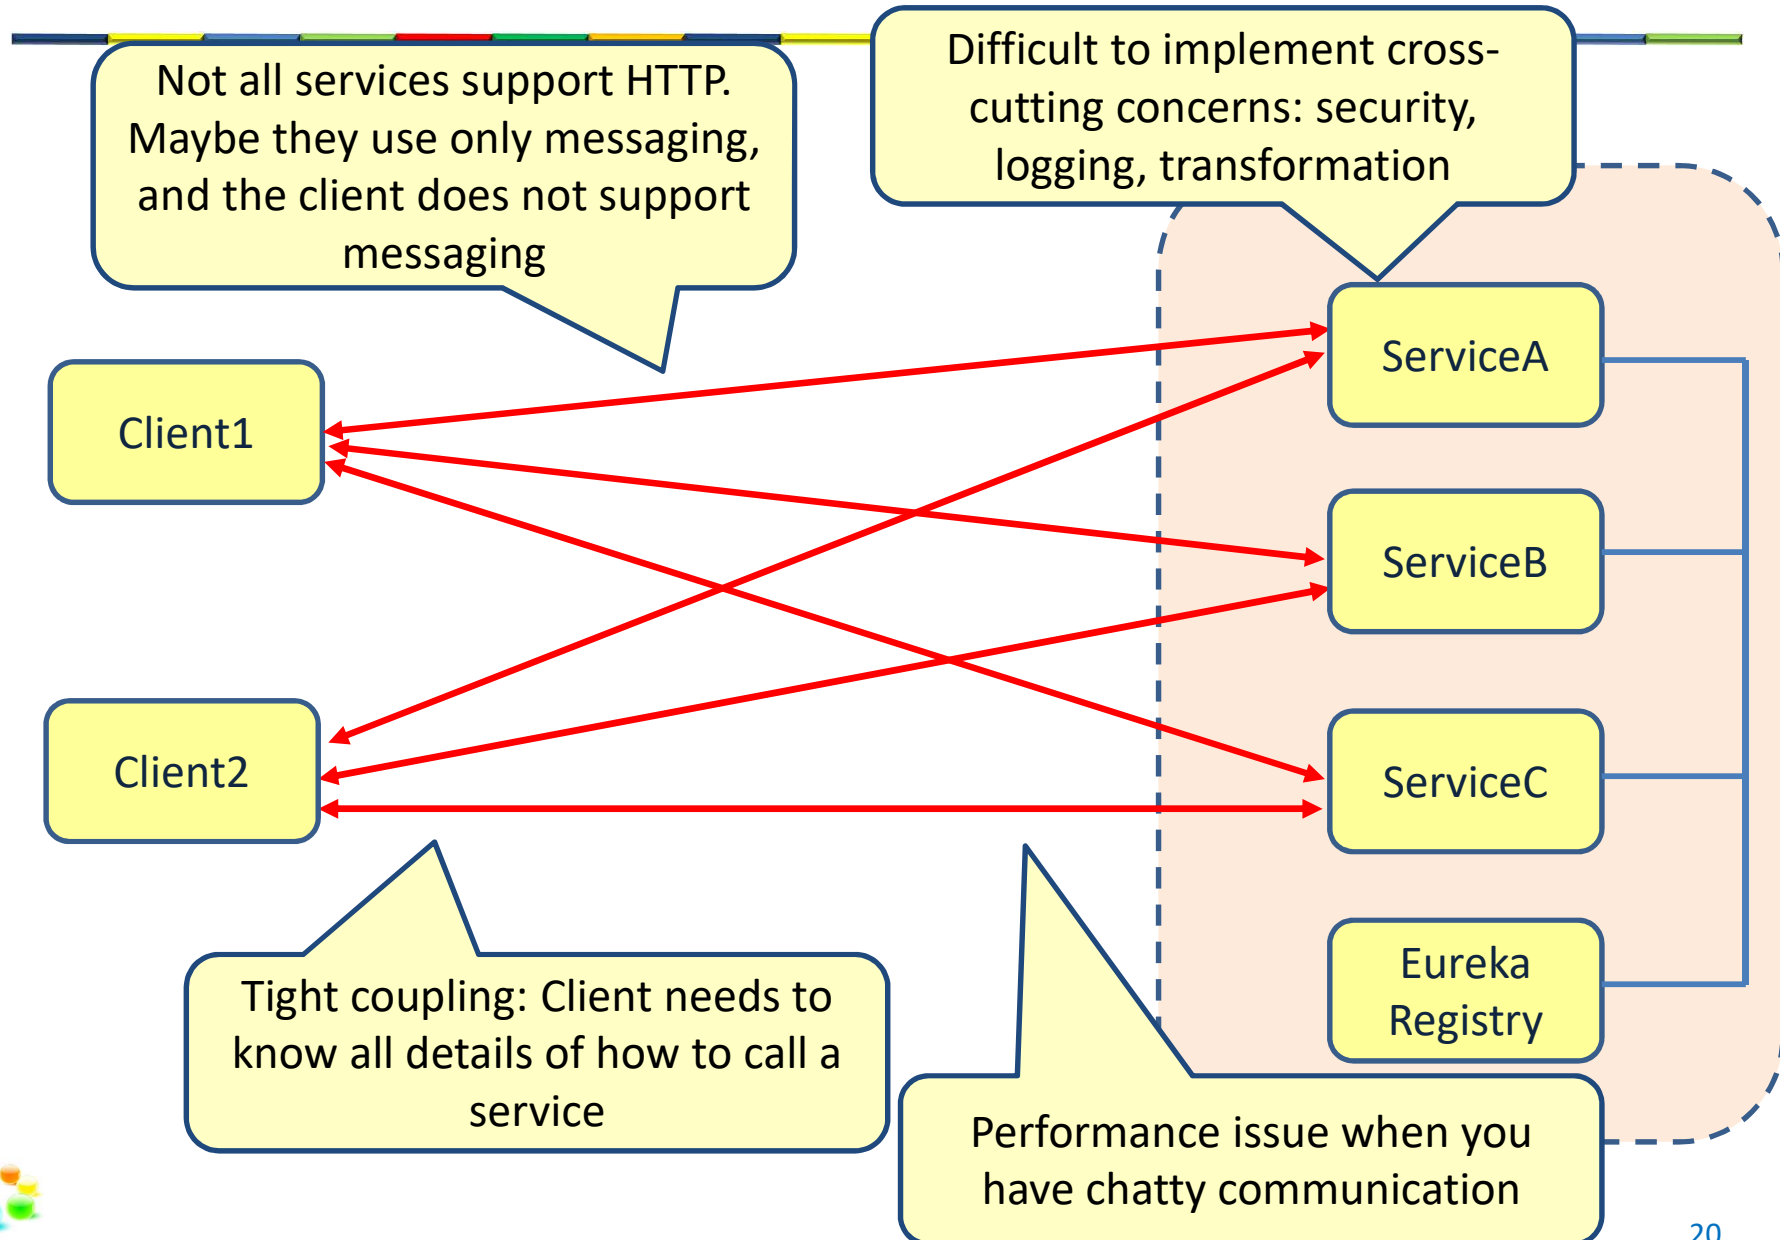


NETFLIX ZUUL

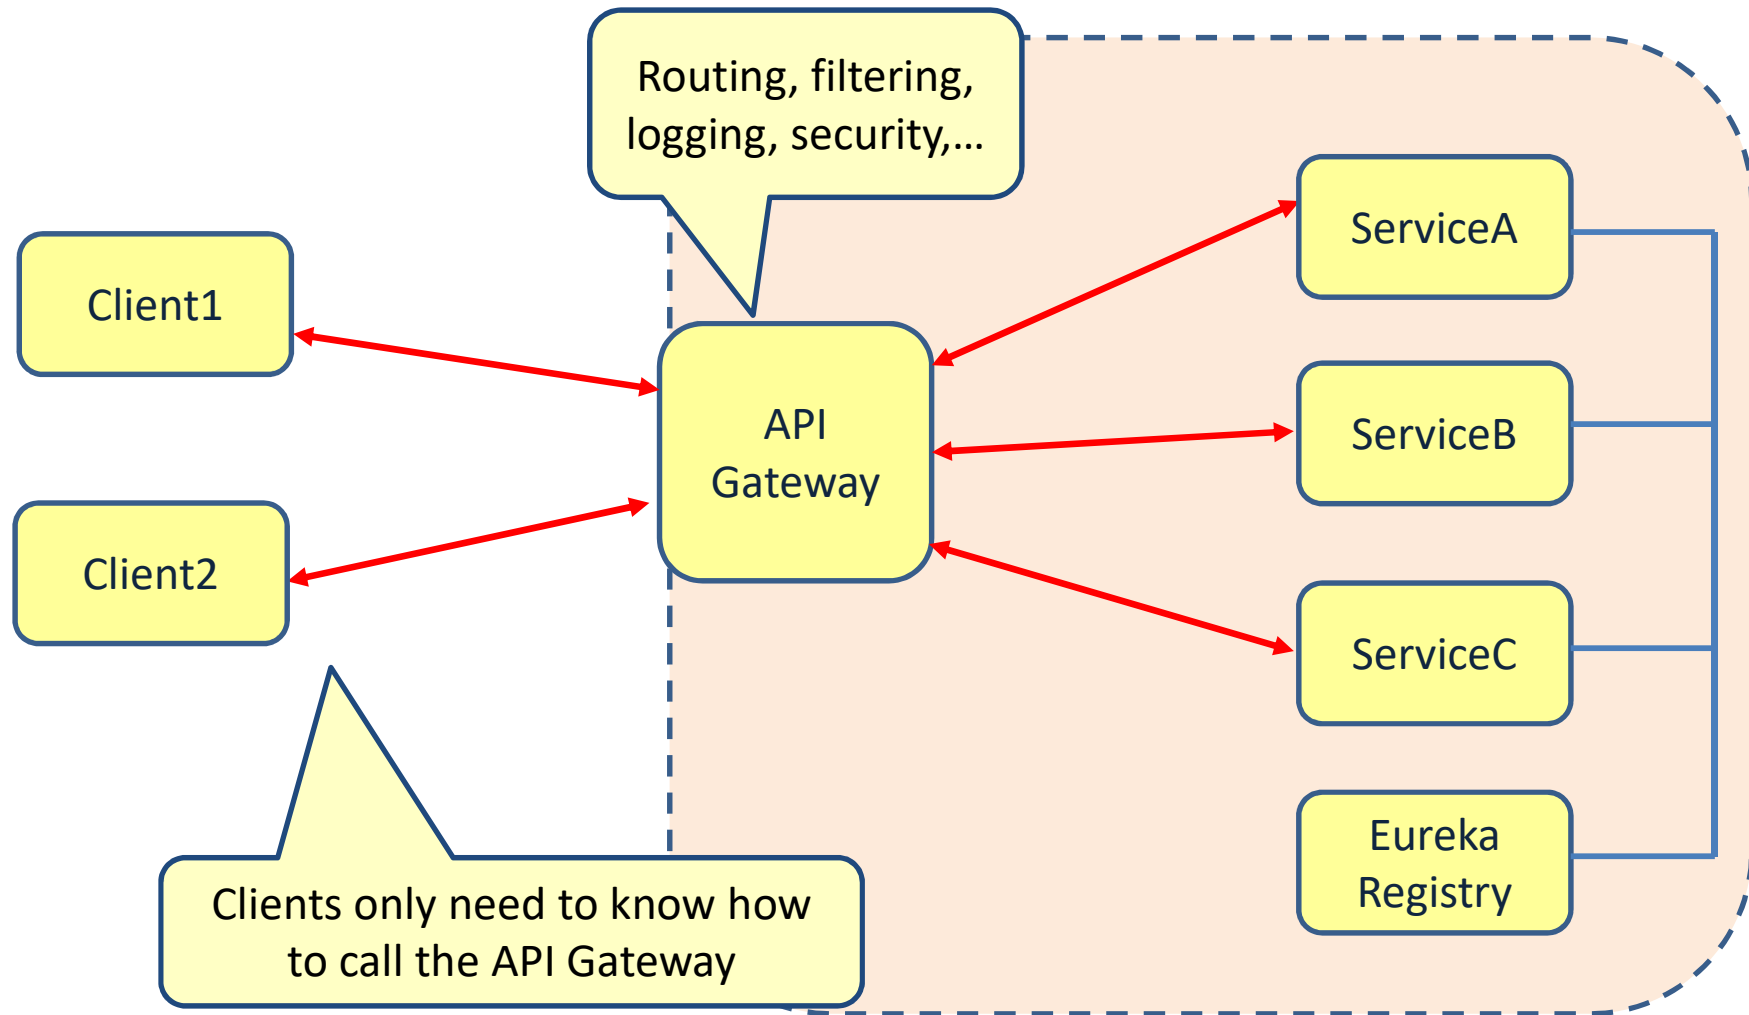
API GATEWAY: ZUUL




Adding clients




Api Gateway



Api Gateway and registry service




API Gateway Filters



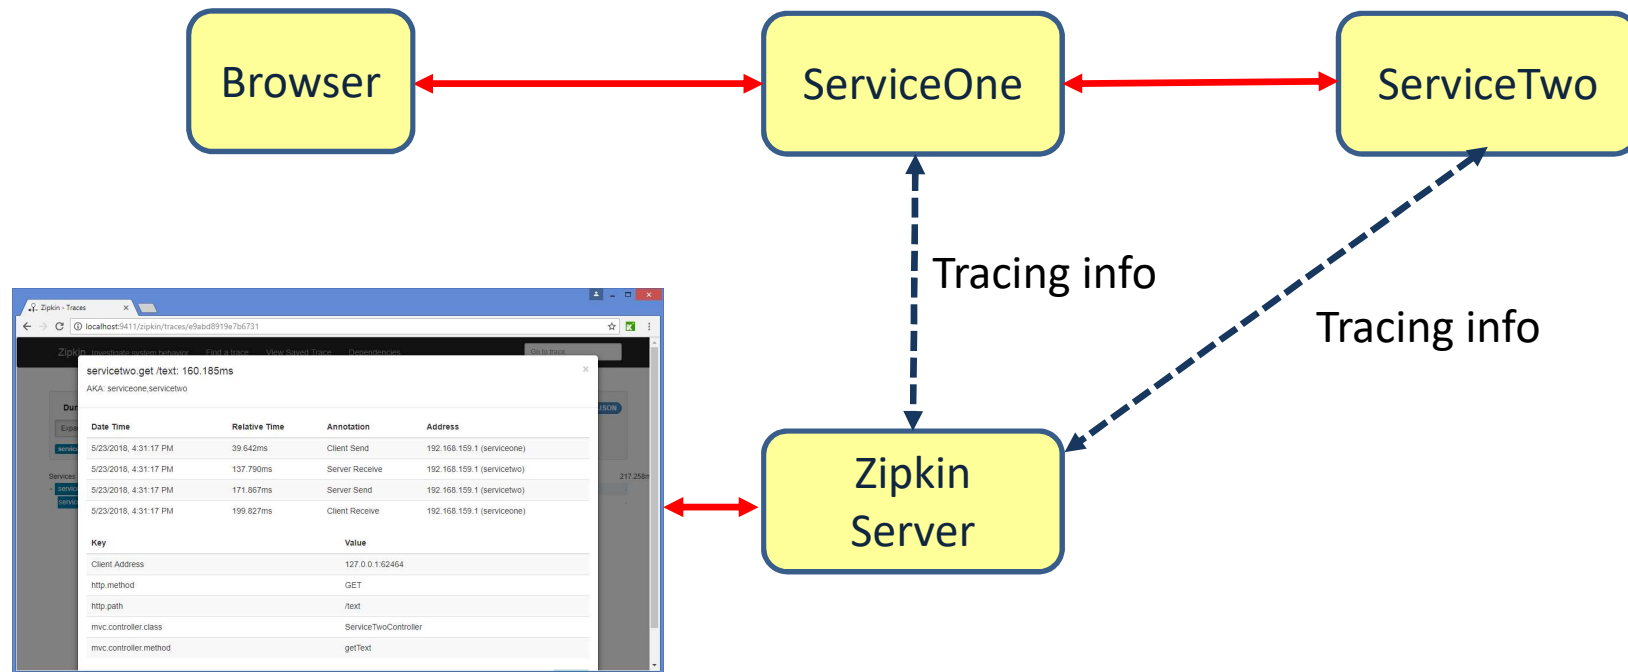
Spring cloud Sleuth

- Span: an individual operation
- Trace: a set of spans

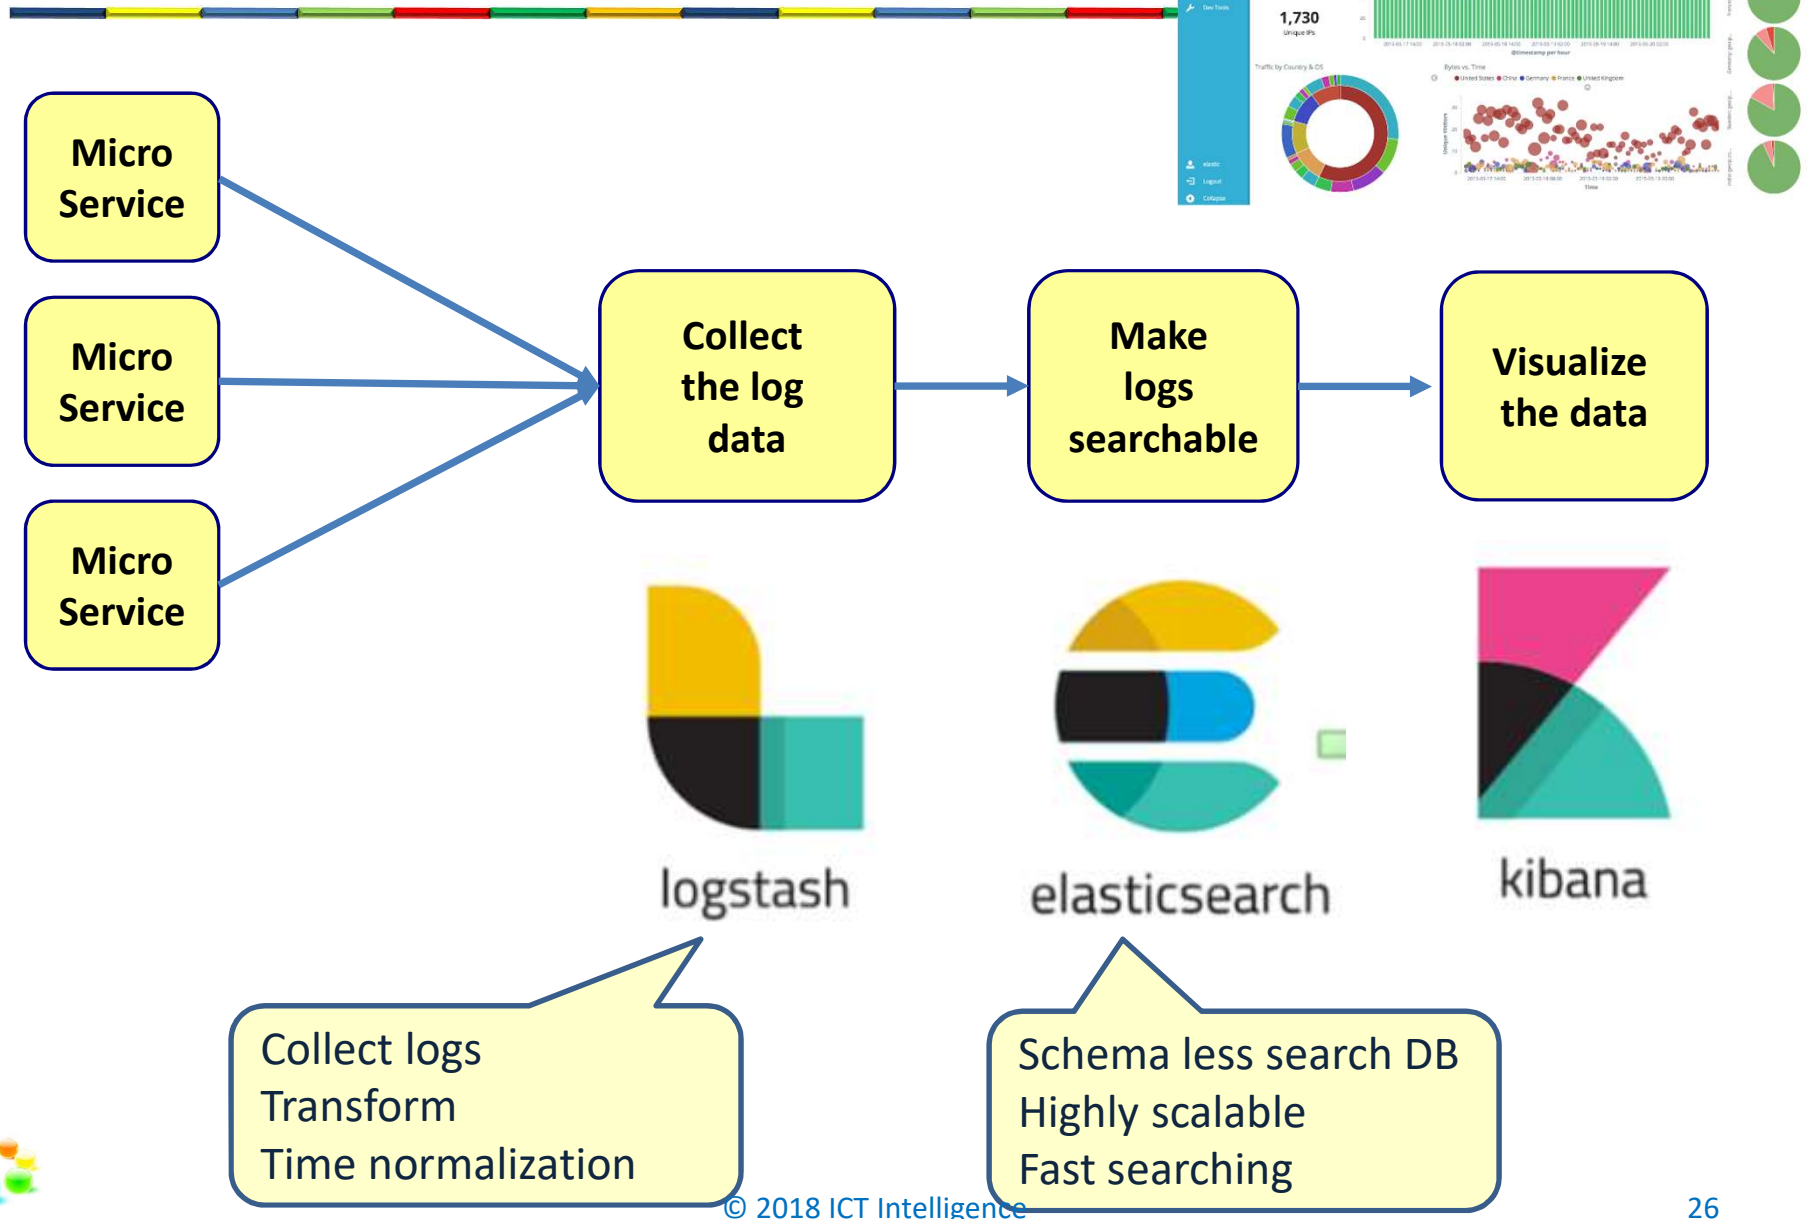


Zipkin

- Centralized tracing server
 - Collects tracing information
- Zipkin console shows the data



ELK stack





NETFLIX

RESILLIENCE: HYSTRIX

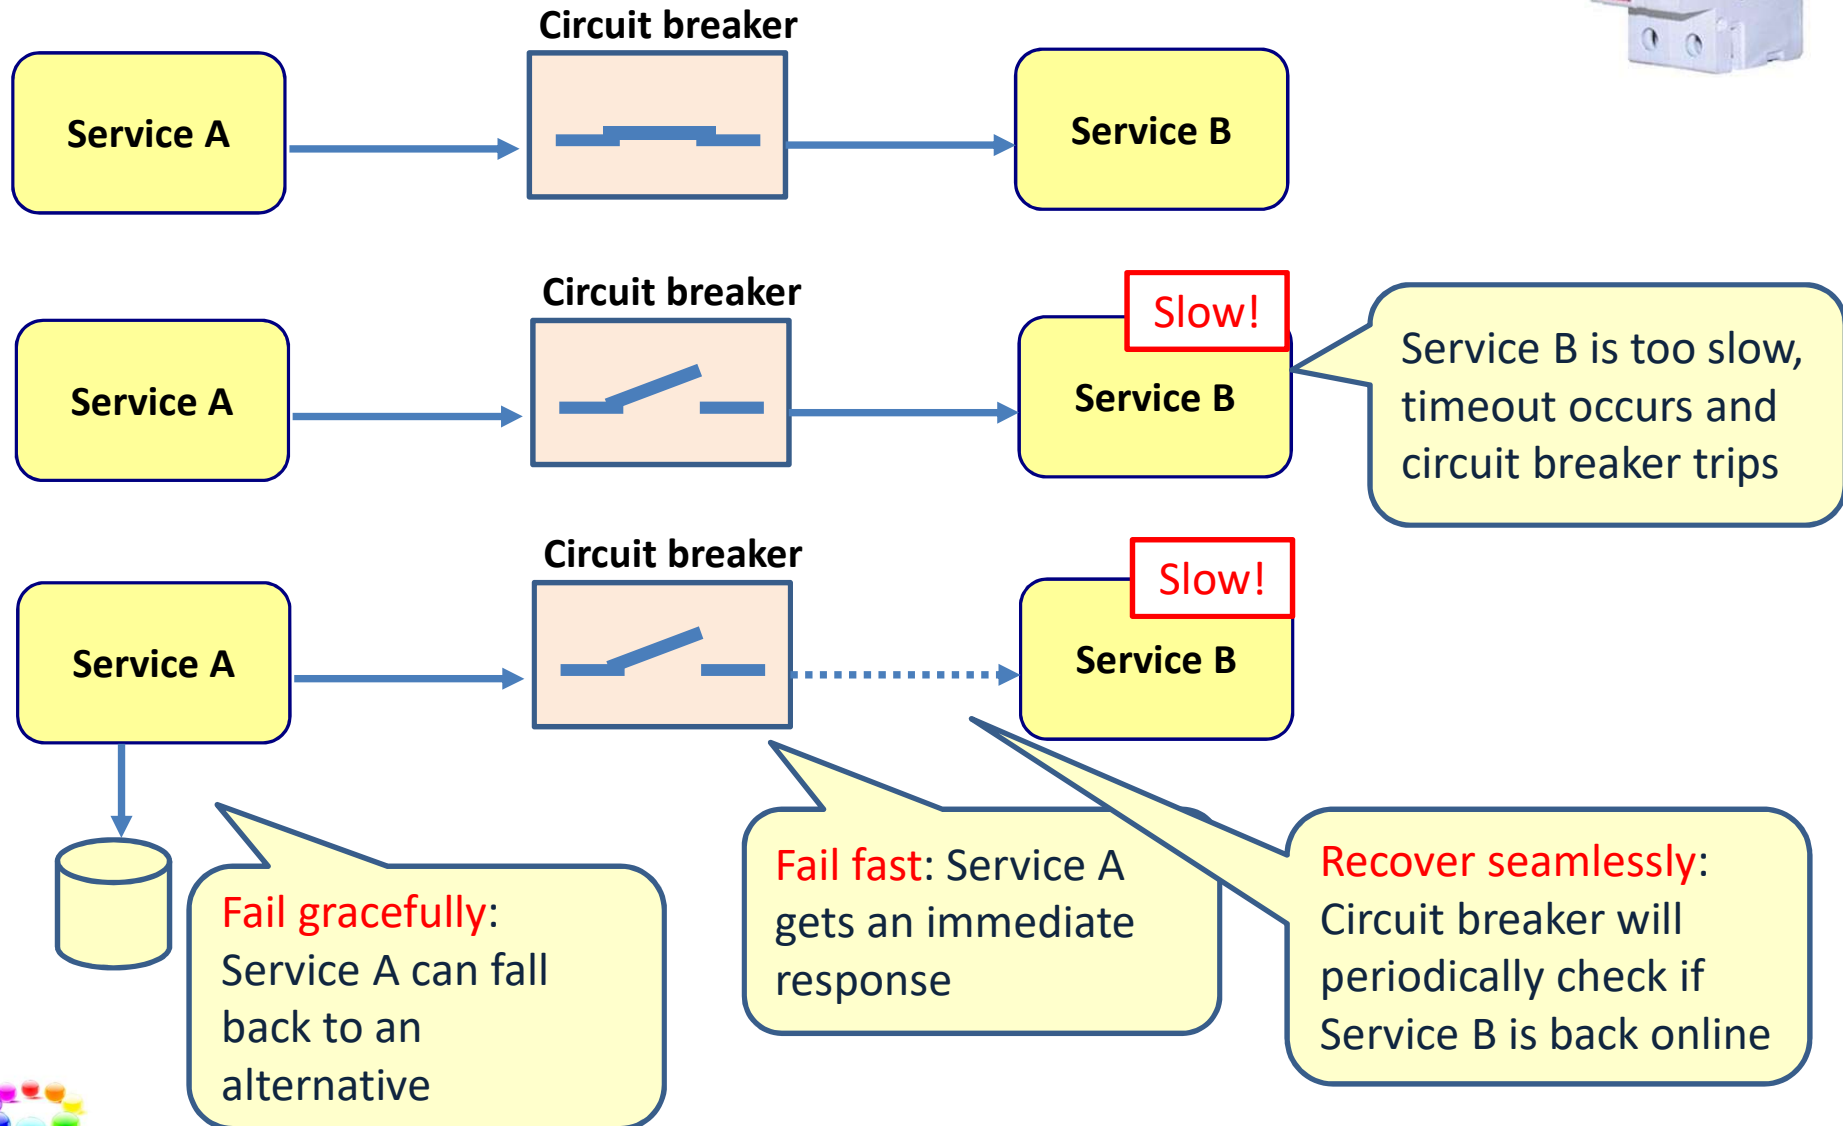


Timeouts

- Put timeouts on all out-of-process calls.
 - Other services
 - Database
 - File system
- Log when timeouts occur
 1. Pick a default timeout
 2. Monitor
 3. Adjust




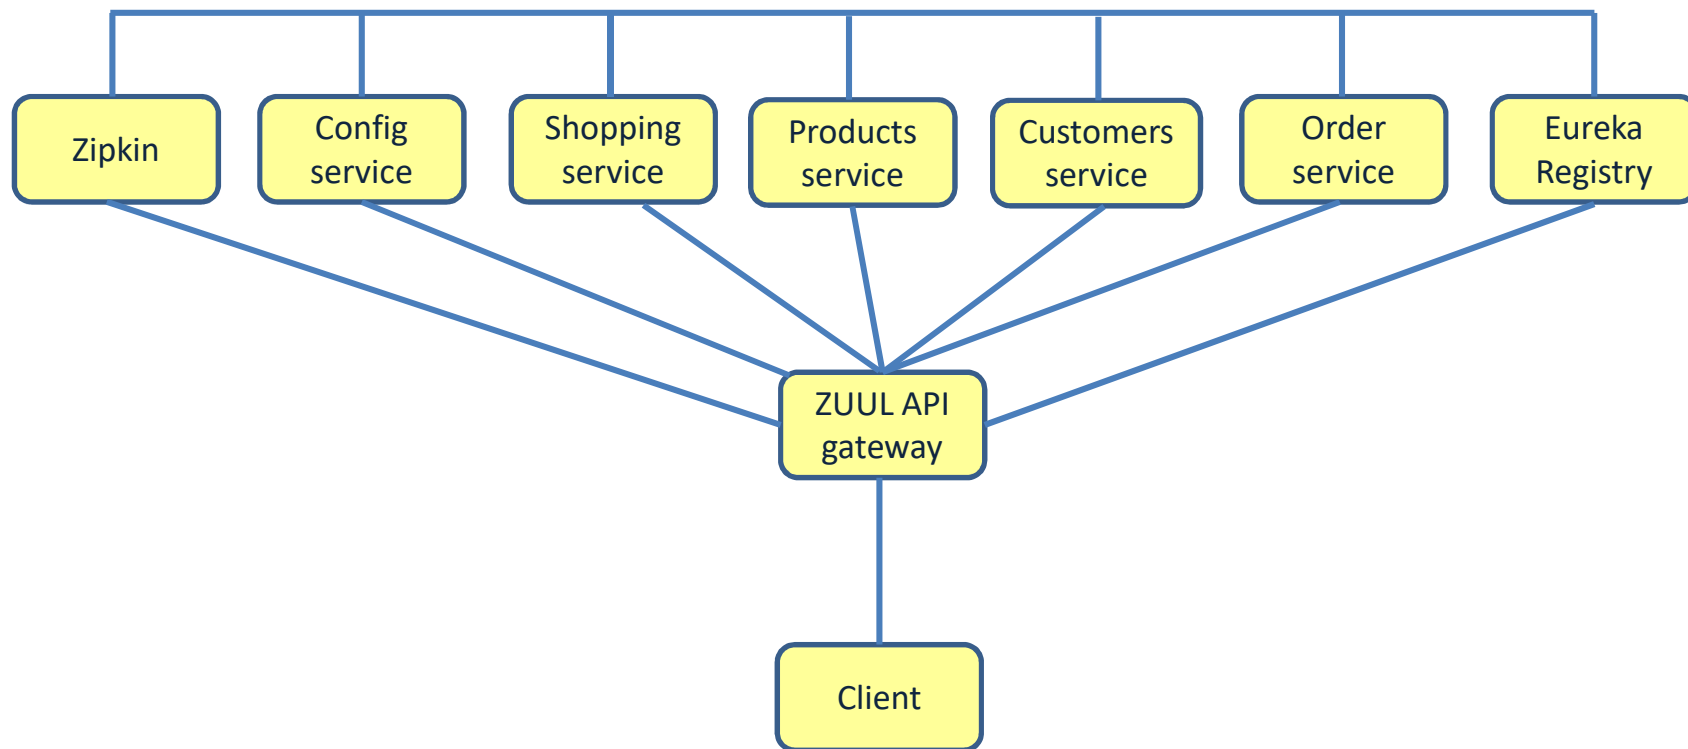
Circuit breaker

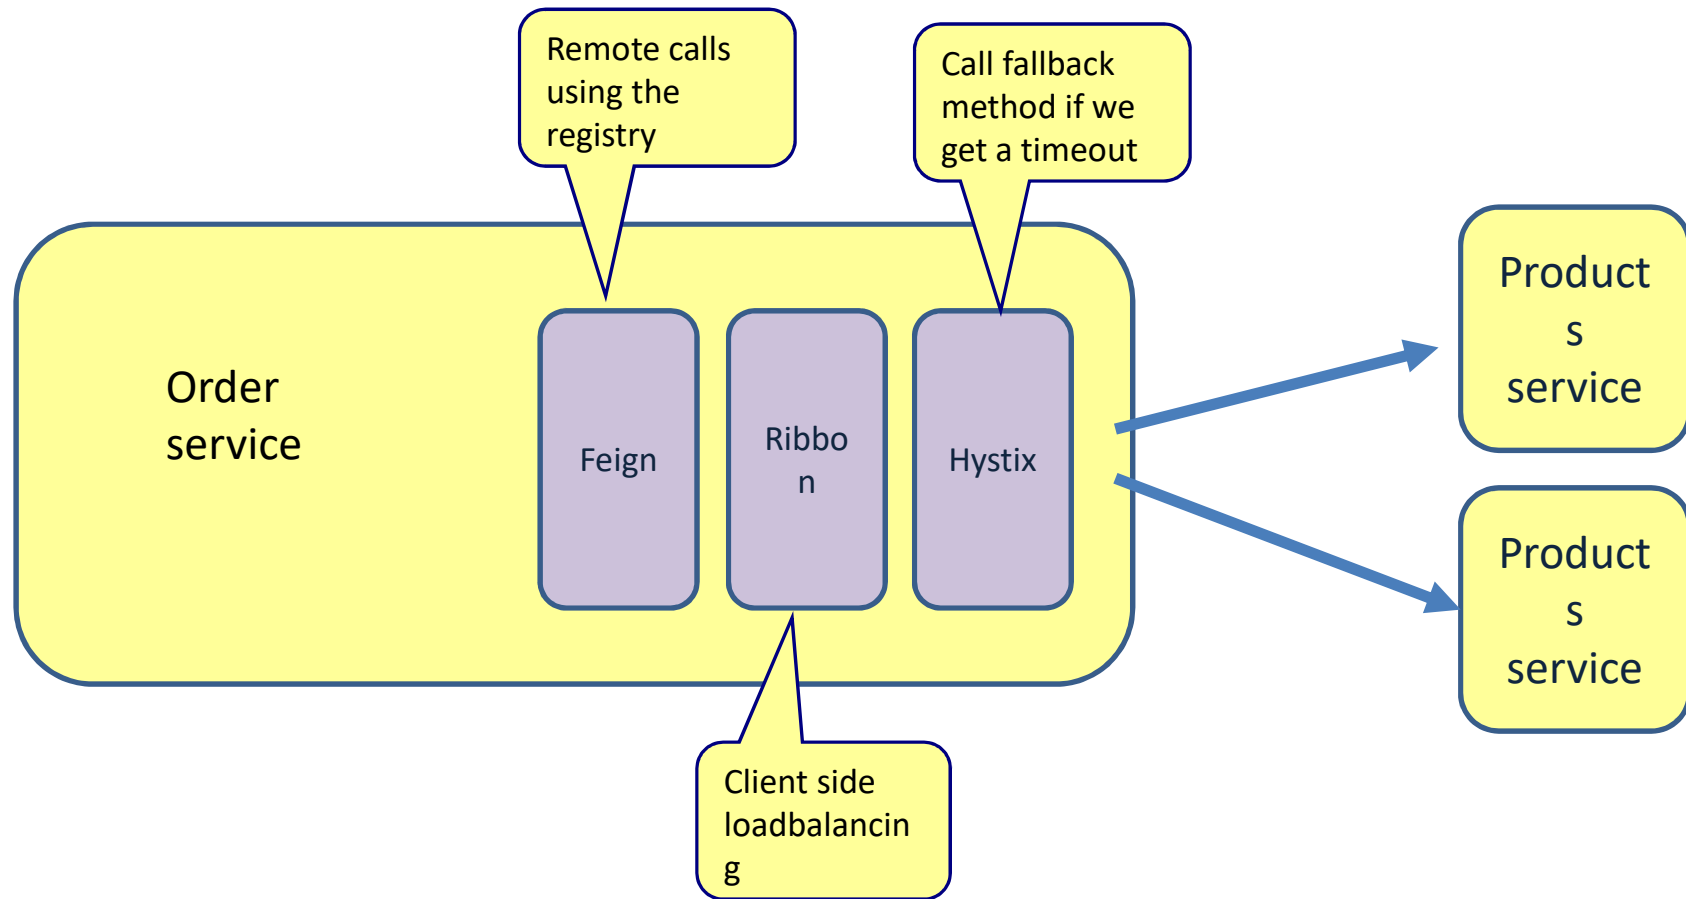


Hystrix bulkheads

- Hystrix uses a common thread pool for all remote calls








CQRS



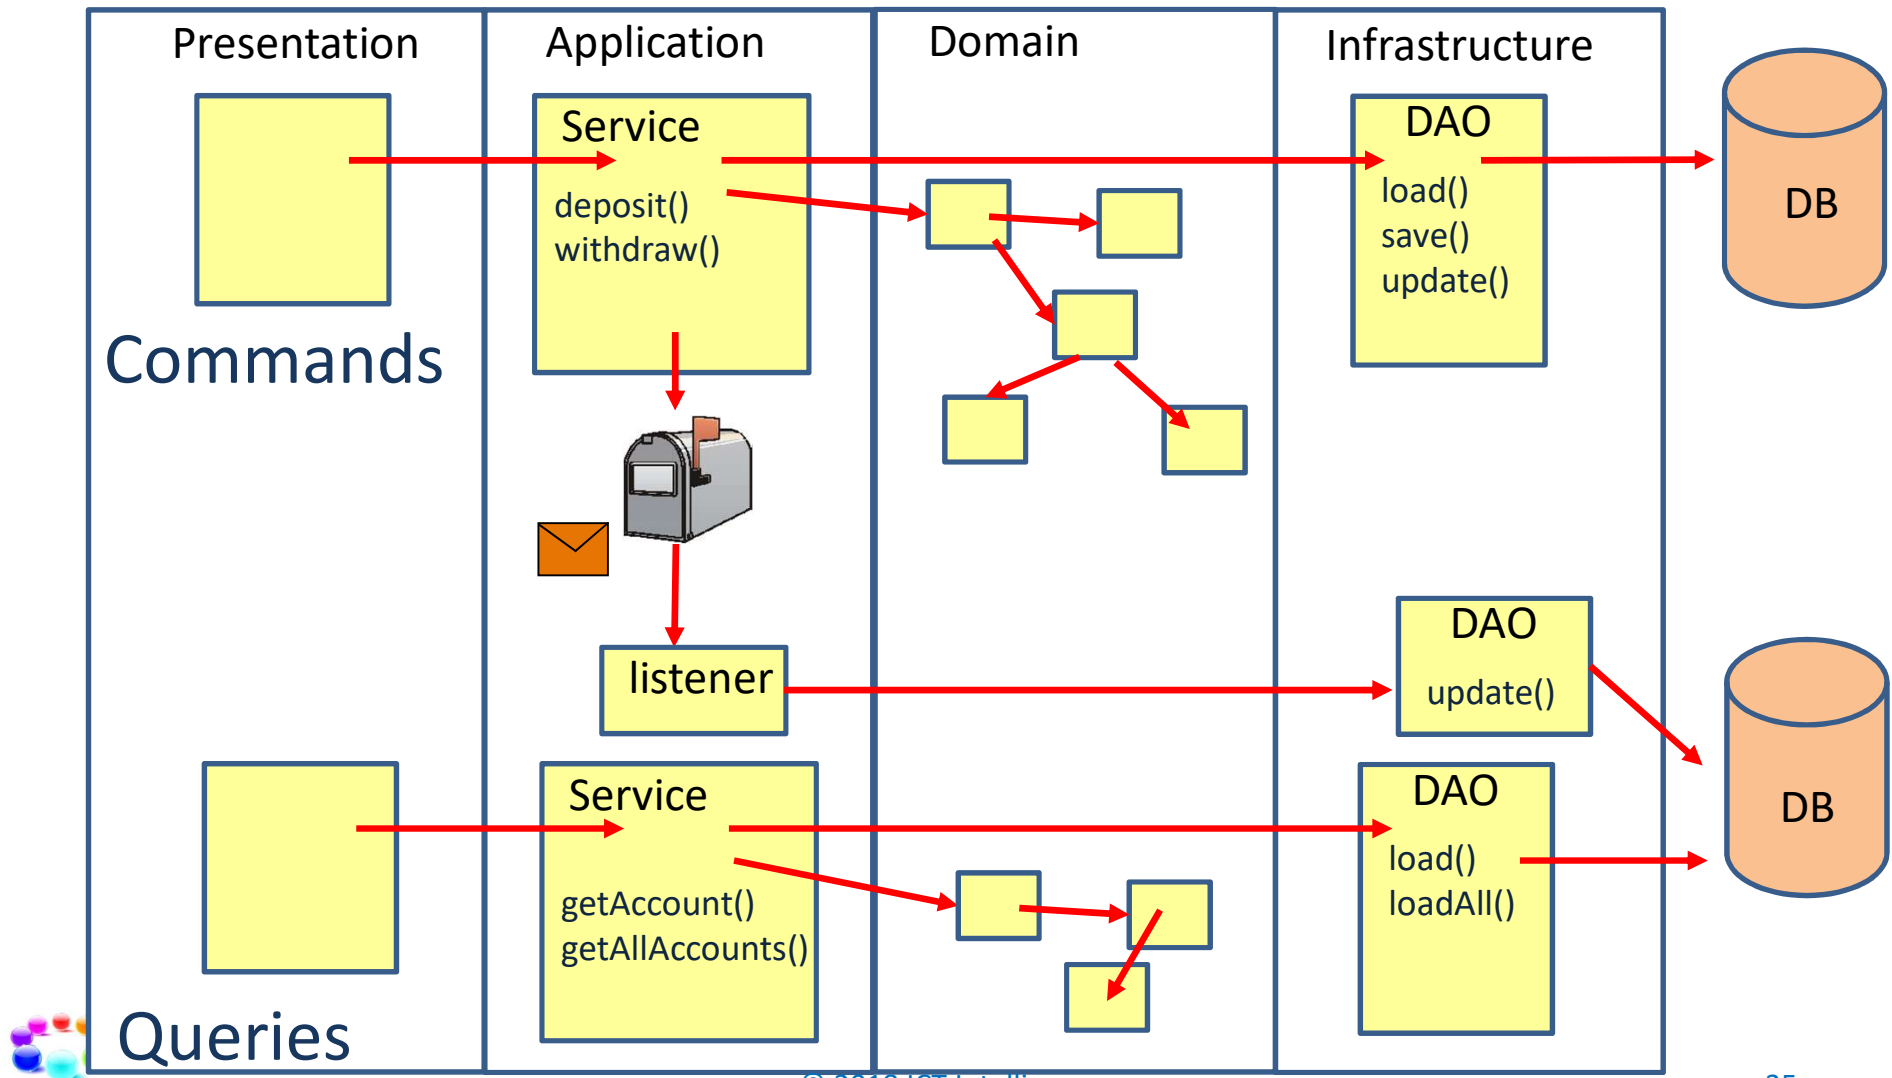
CQRS

- Domain model is used for commands
- View model is used for queries



Eventual consistency

- Views will become eventual consistent




REACTIVE REST WITH SPRING WEBFLUX




Imperative versus reactive

- Synchronous, blocking




- Asynchronous, non-blocking




Reactive Web

- Synchronous, blocking



- Asynchronous, non-blocking

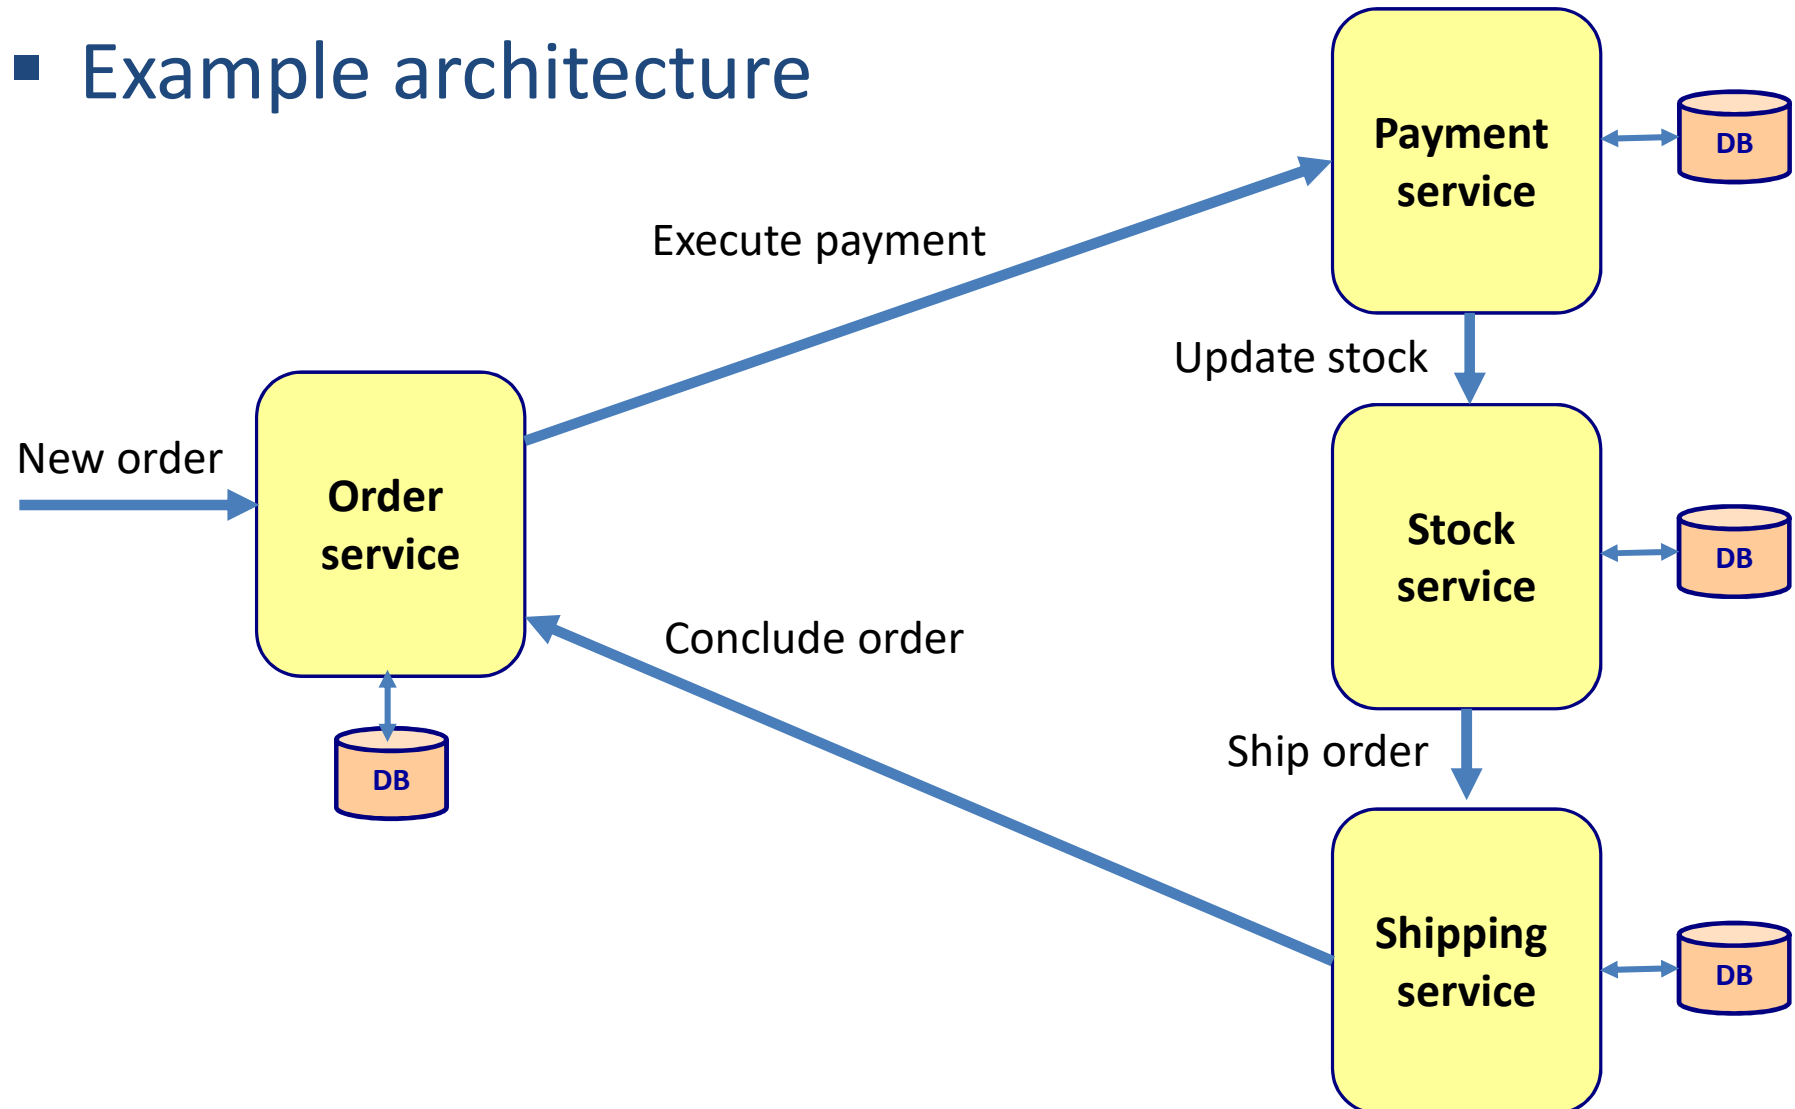


TRANSACTIONS

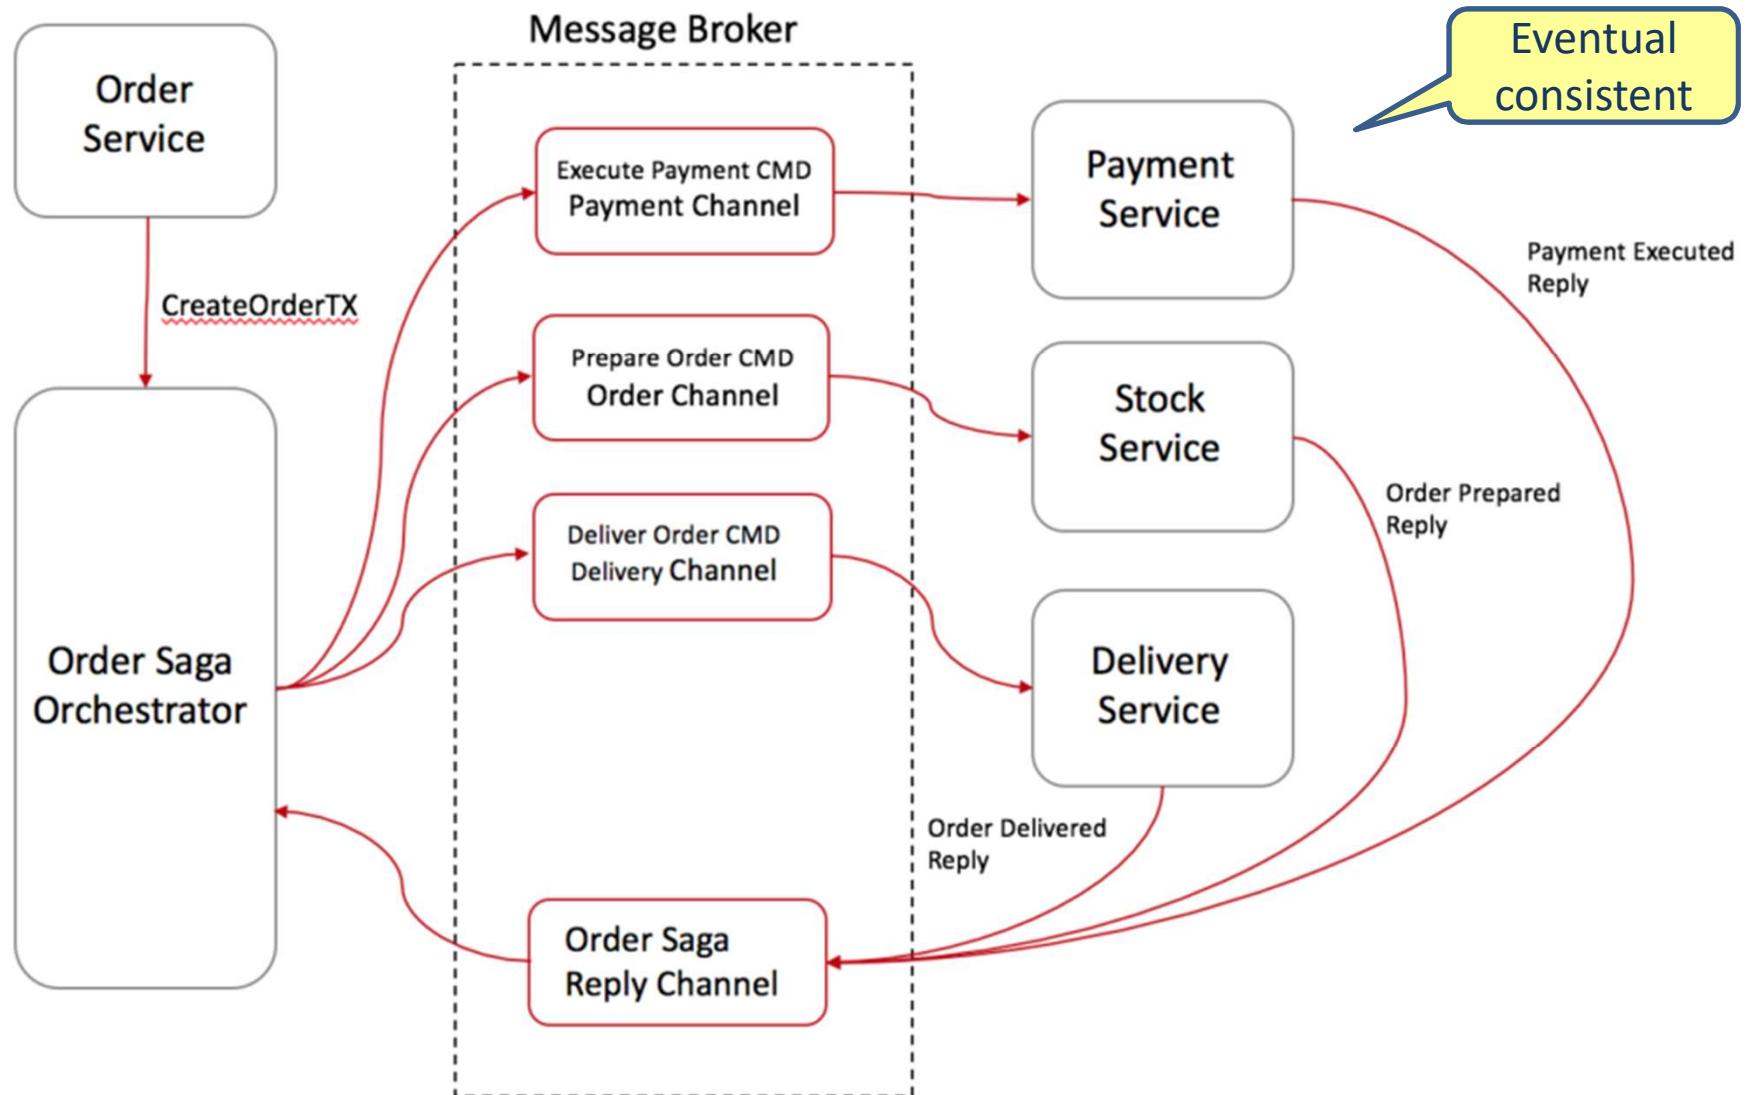


Saga pattern


- Example architecture



Saga with command/orchestration



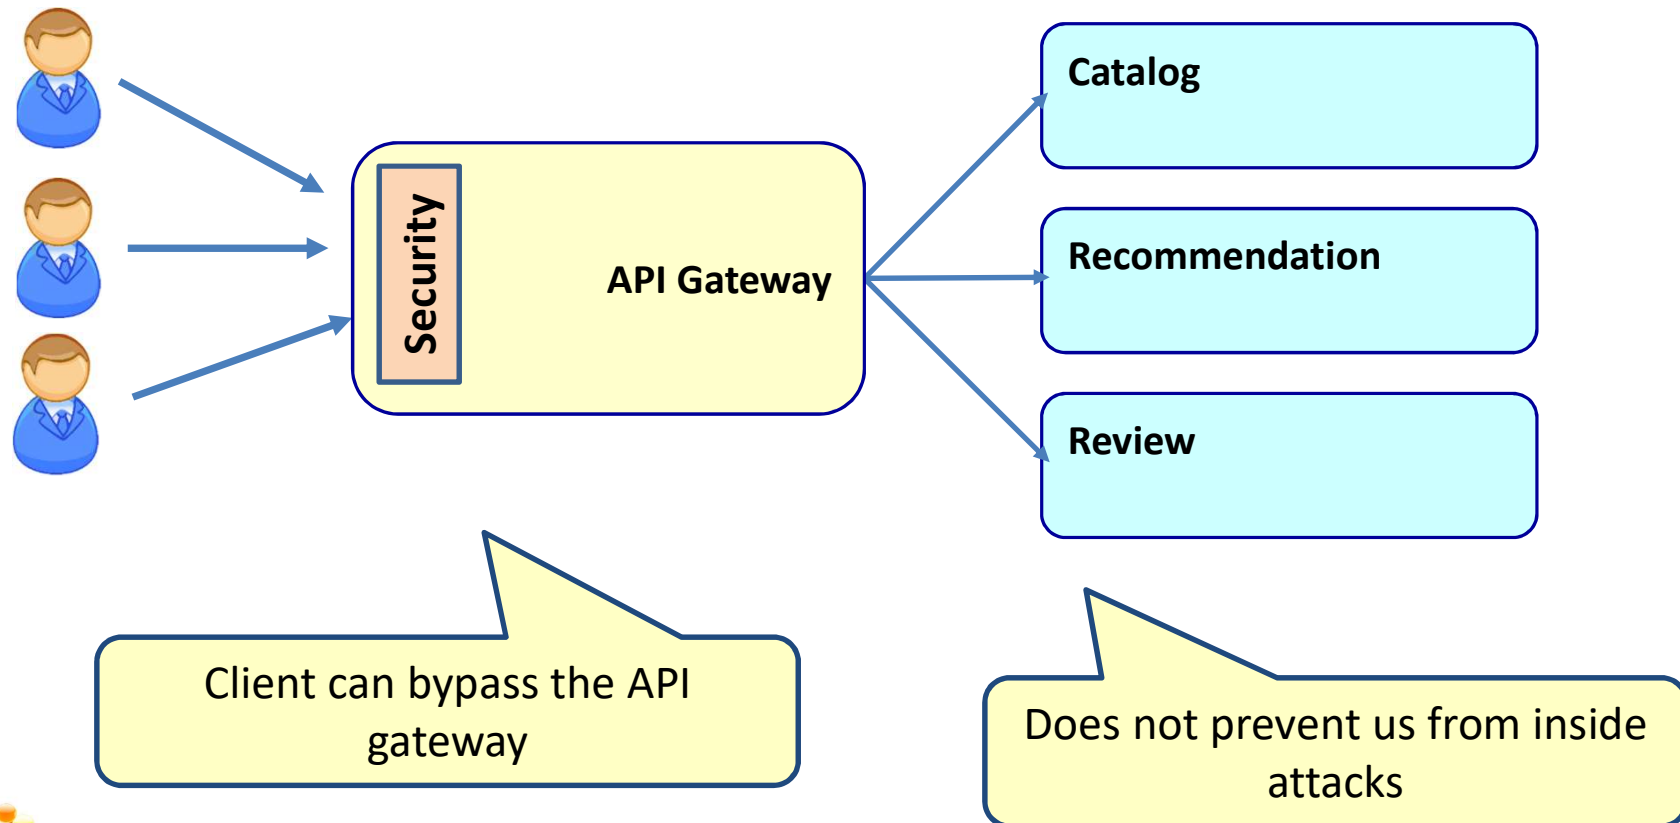
Saga with events/choreography



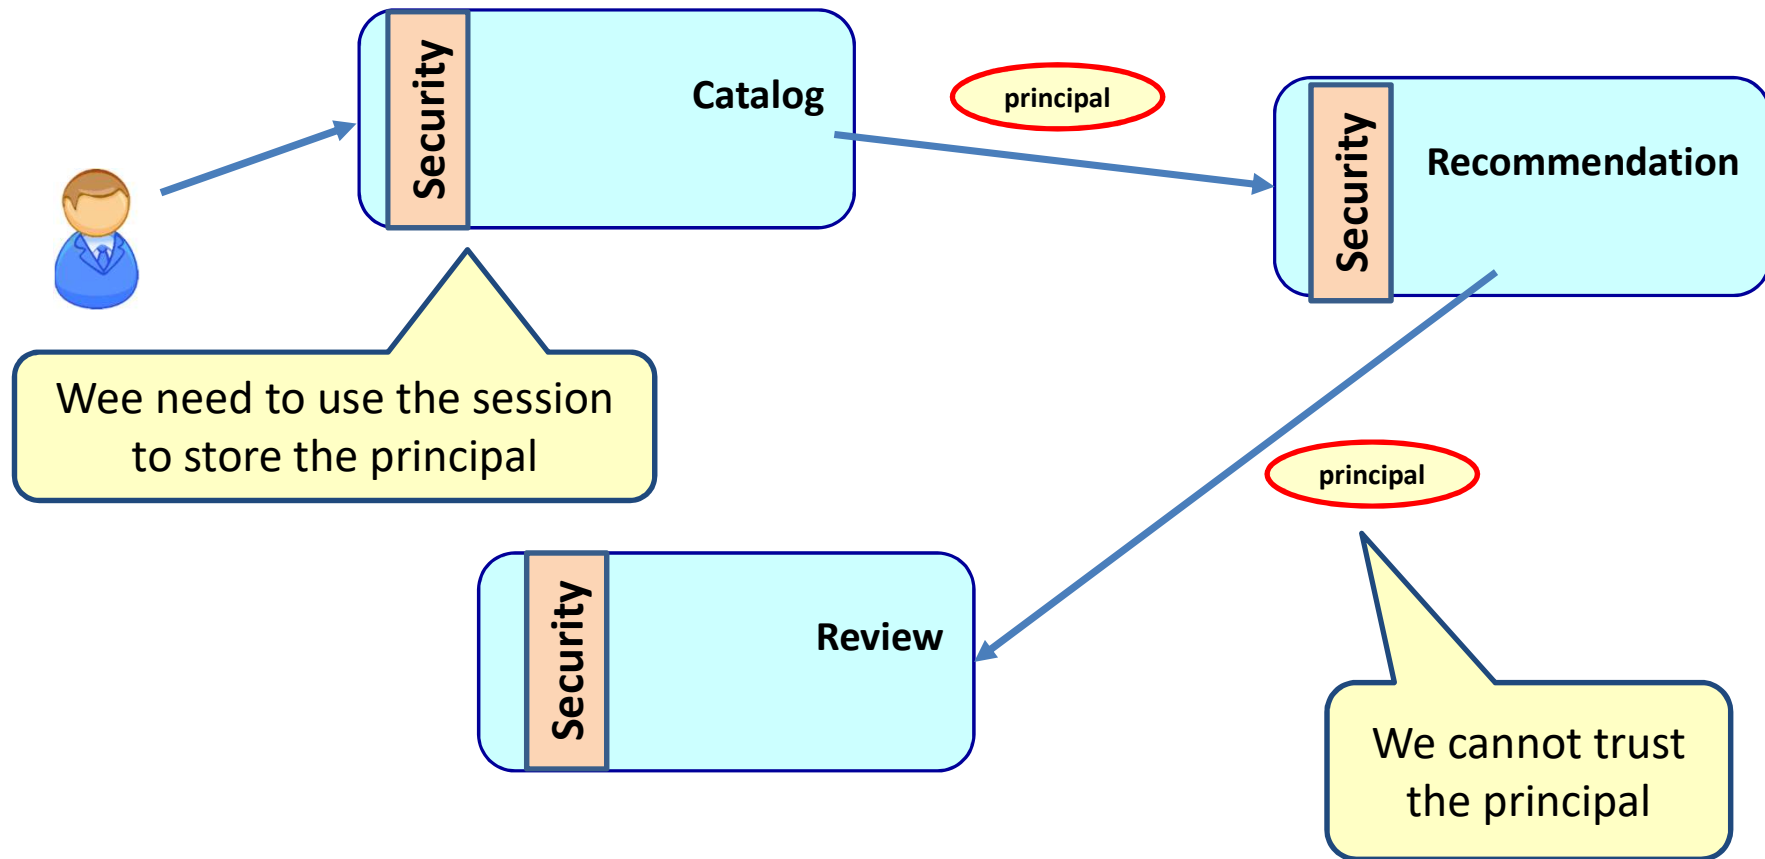
SECURE THE MICROSERVICE ARCHITECTURE



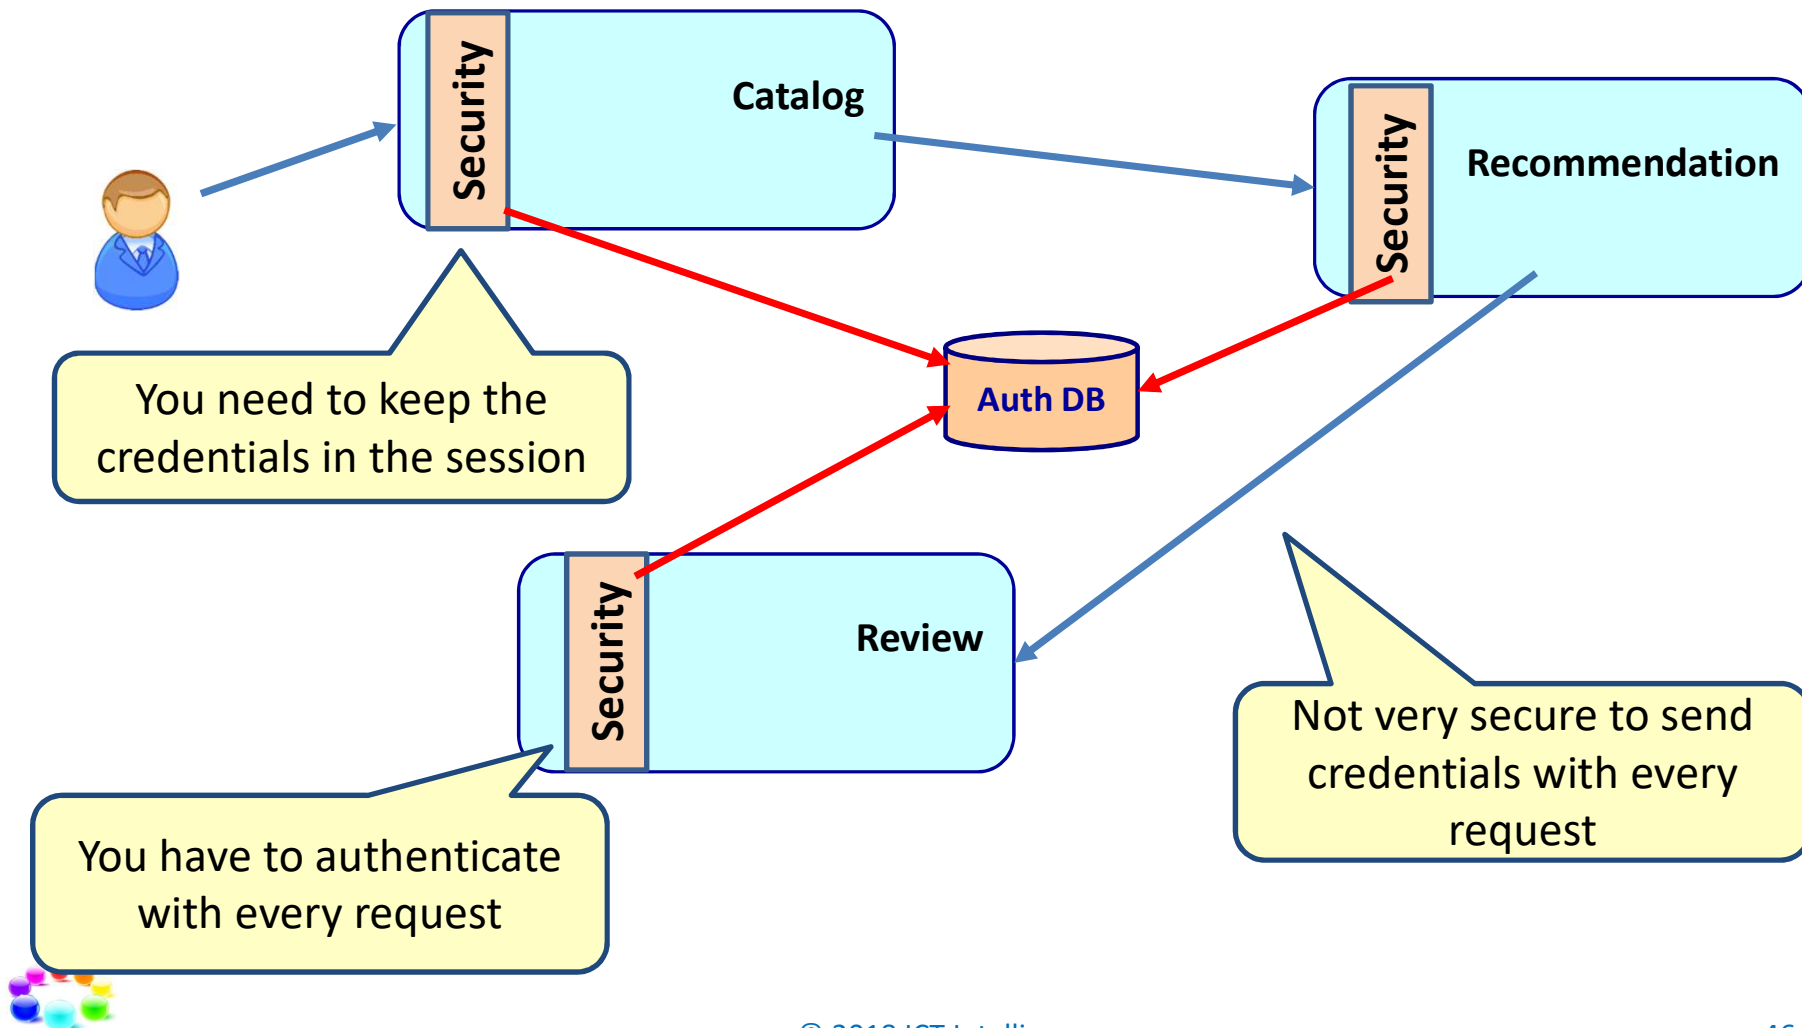
Secure the API gateway




Send principal with every request



Send userid/password with every request

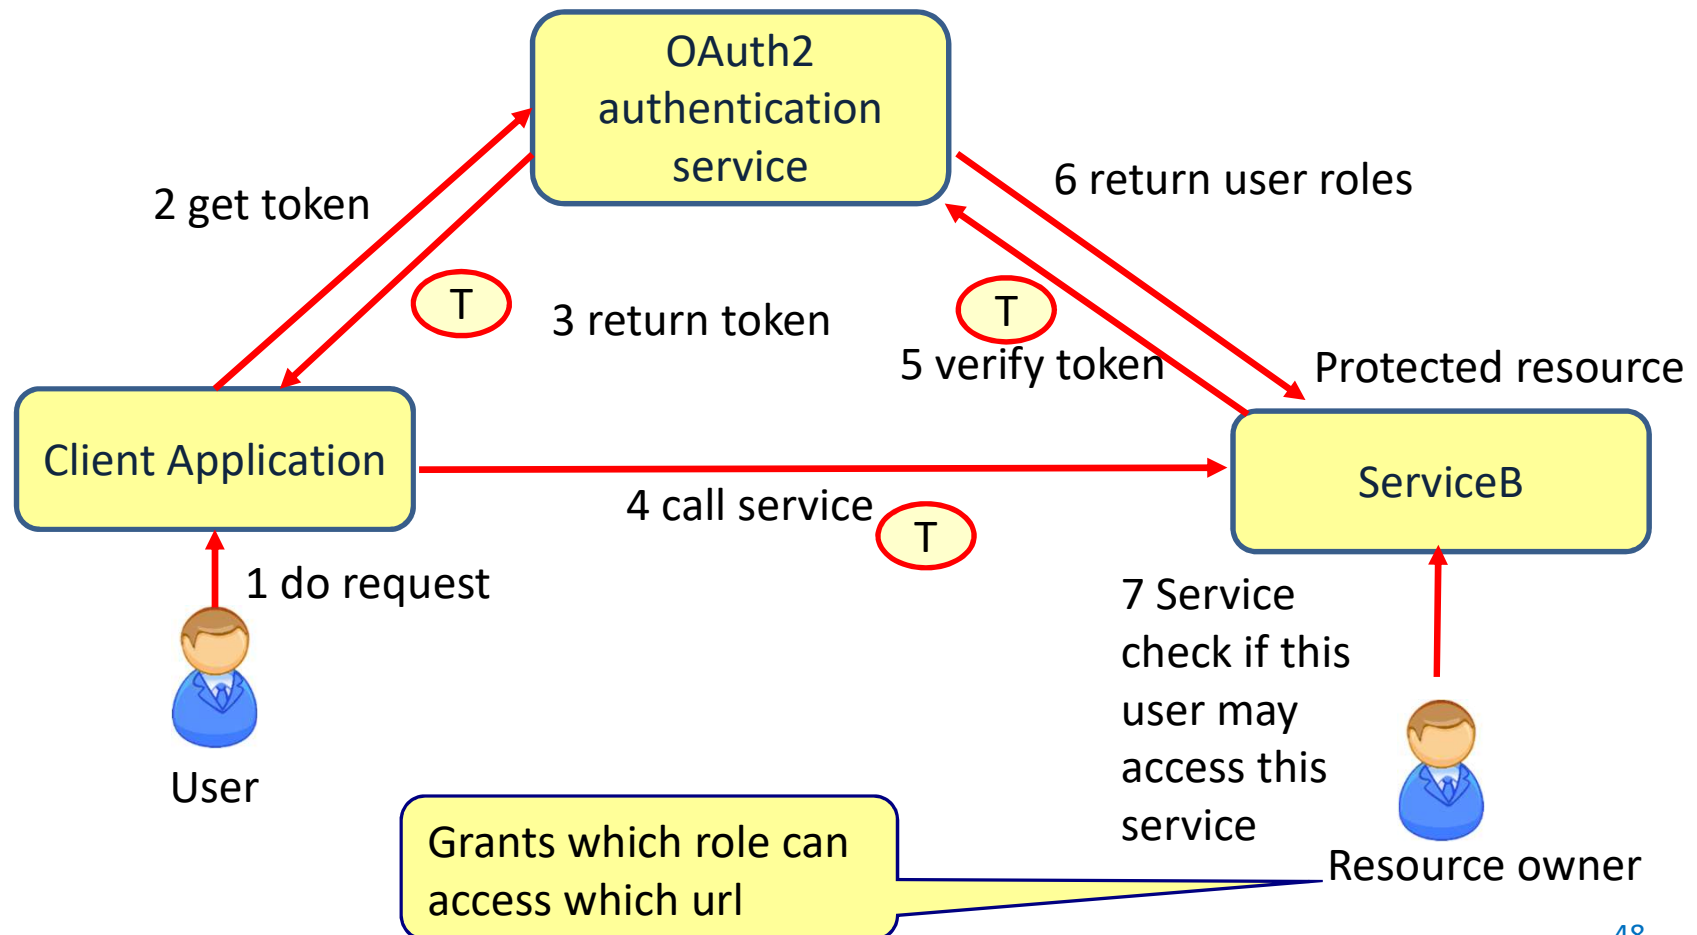


OAuth2: Token based security

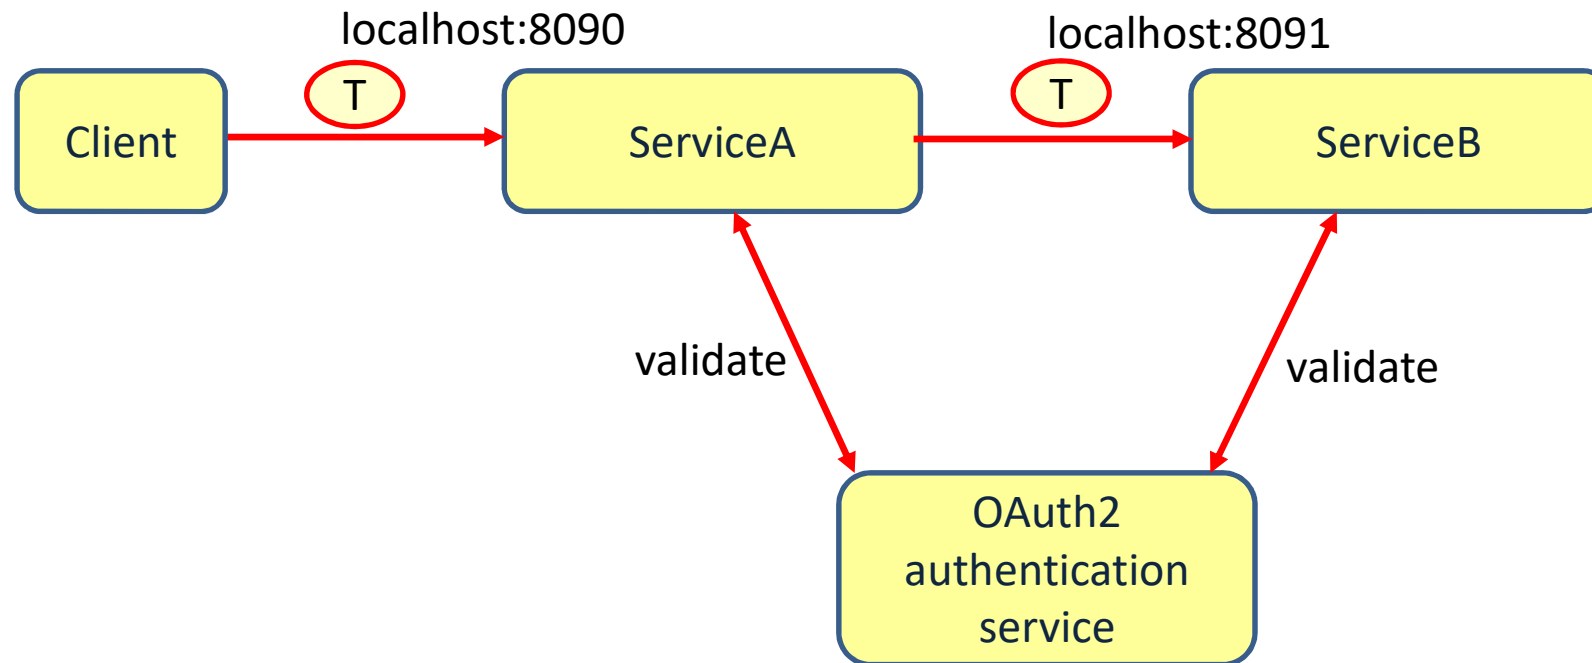


How does OAuth2 work

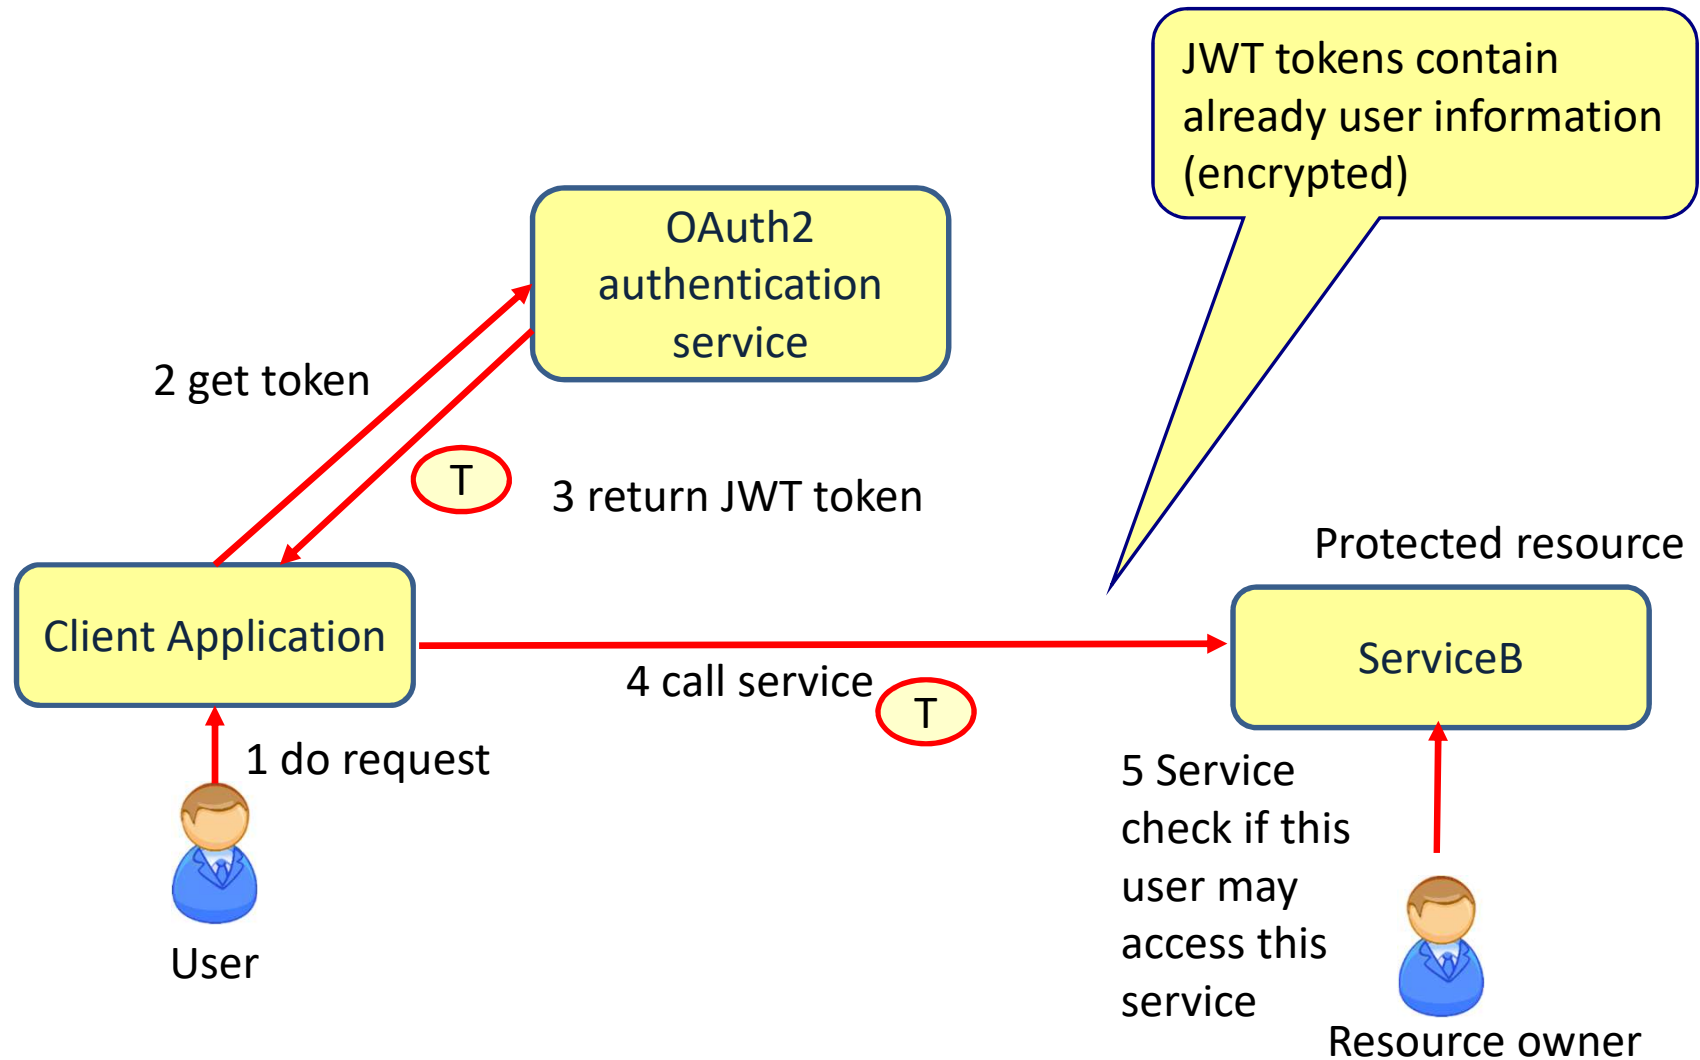
- Token based authentication and authorization framework



Propagate the token



JWT tokens



EVENT DRIVEN ARCHITECTURE





Event driven architecture

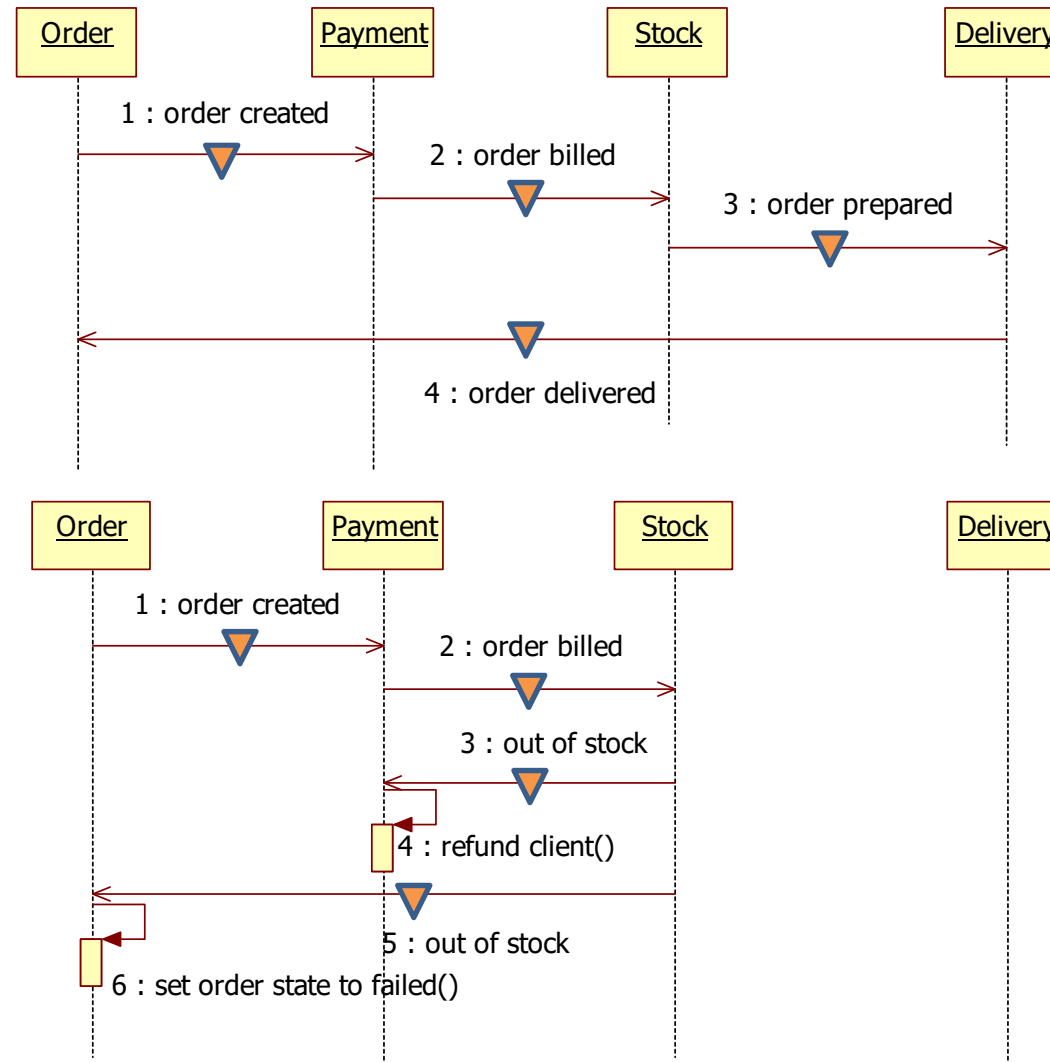
- Loosely coupled services
- Asynchronous
 - No blocking calls
 - No threads that are just waiting
- Flexible
 - Publish-subscribe
 - Easy to add new publishers
 - Easy to add new listeners
- Buffer
 - If a service is slow or down




Synchronous (REST) calls



Asynchronous events (messaging)




Problems with traditional messaging middleware



- If the consumer is temporally not available (or very slow) the message middleware has to store the messages
 - This restricts the volume of messages and the size of the messages
 - Eventually the message broker will fail



Problems with traditional messaging middleware



- If the consumer has a bug, and handles the messages incorrectly, then the messages are gone.
 - Not fault-tolerant



Apache Kafka



- Created by Linked In



- Characteristics

- High throughput

- Distributed

- Unlimited scalable

- Fault-tolerant

- Reliable and durable

- Loosely coupled Producers and Consumers

- Flexible publish-subscribe semantics

High Volume:


- Over 1.4 trillion messages per day
- 175 terabytes per day

High Velocity:

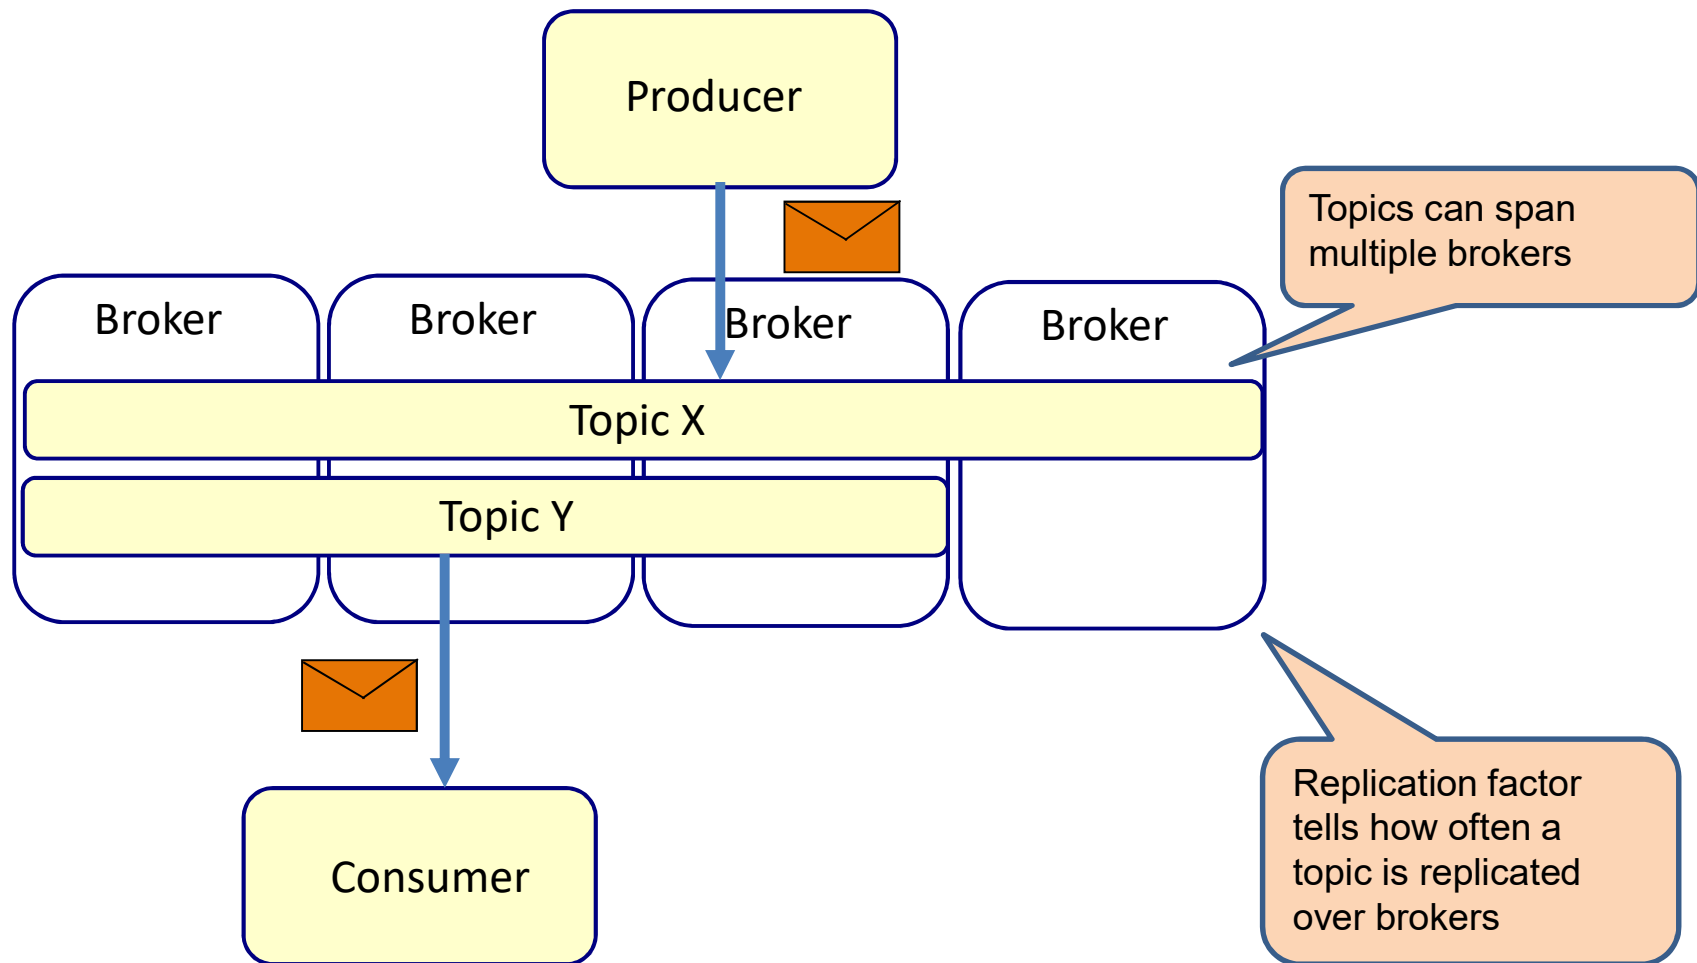
- Peak 13 million messages per second
- 2.75 gigabytes per second



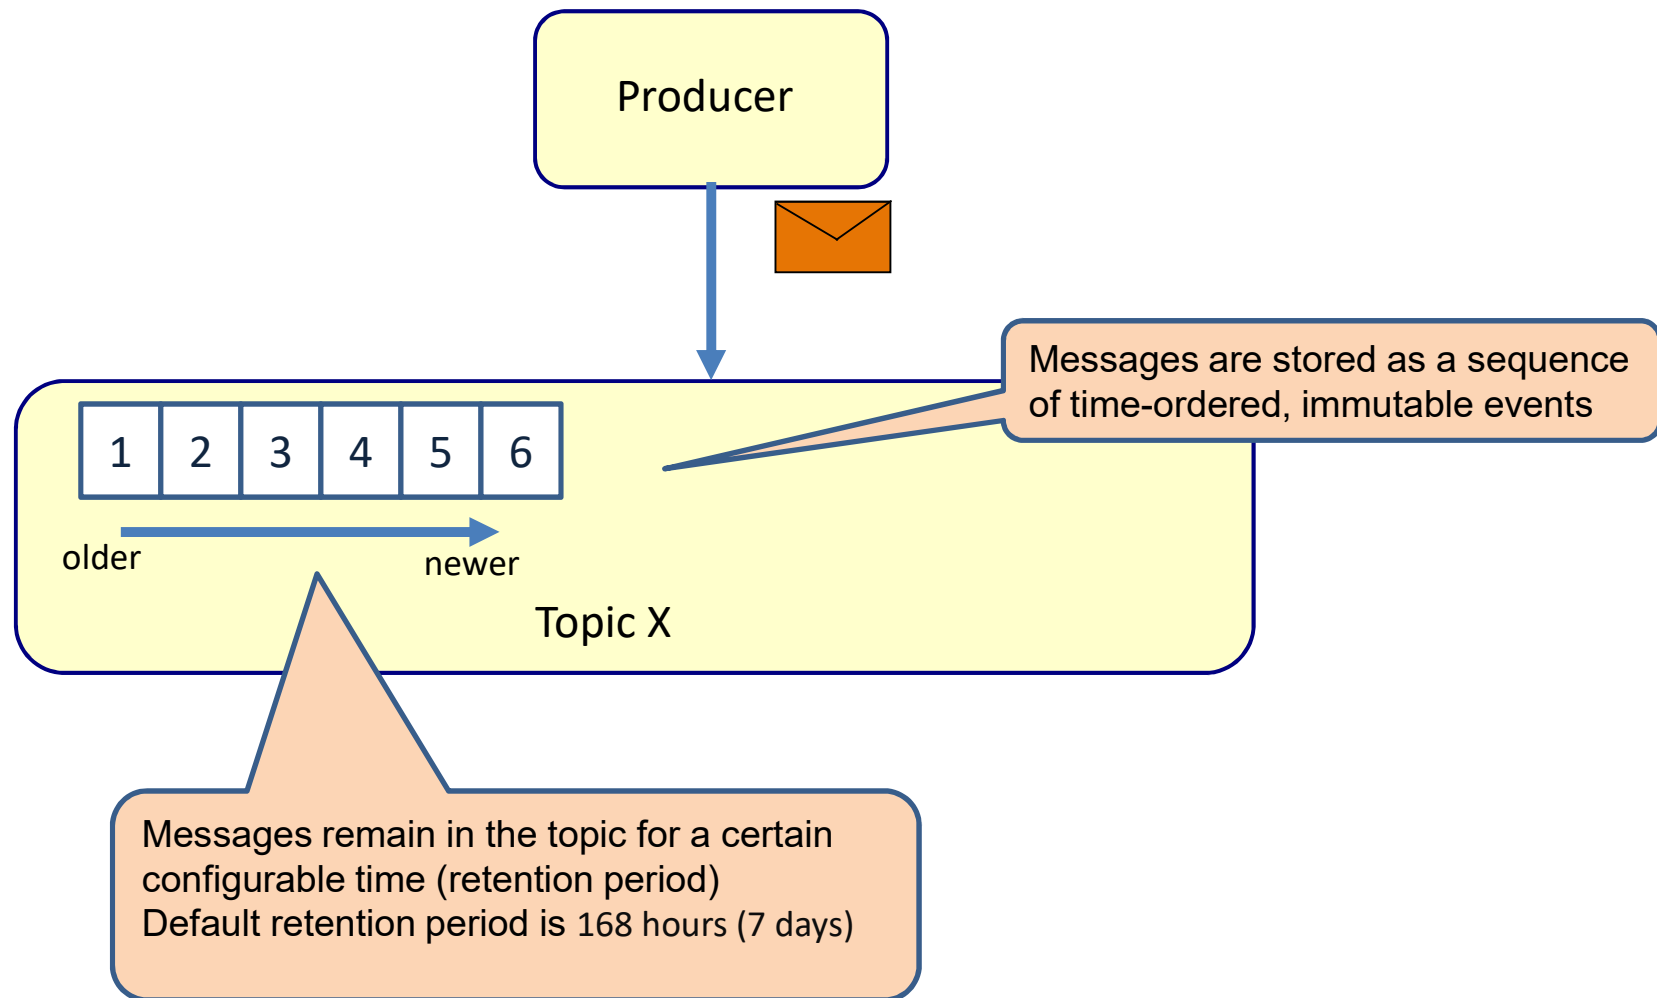
Cluster of Brokers




Topics



Event sourcing



Scale out partitions




STREAM BASED ARCHITECTURE



Blackboard

- Common data structure
 - Extension is no problem
 - Change is difficult
- Easy to add new components
- Tight coupling for data structure
- Loose coupling for
 - Location
 - Time
 - Technology(?)
- Synchronisation issues




Event sourcing

- Instead of storing the state of an entity in a database, you store the series of events that lead up to the state.
- Storing all of the events increases the analytical capabilities of a business.
- Instead of just asking what the current state of an entity is, a business can ask what the state was at any time in the past




Batch processing

- First store the data in the database
- Then do queries (map-reduce) on the data
- Queries over all or most of the data in the dataset.
- Latencies in minutes to hours

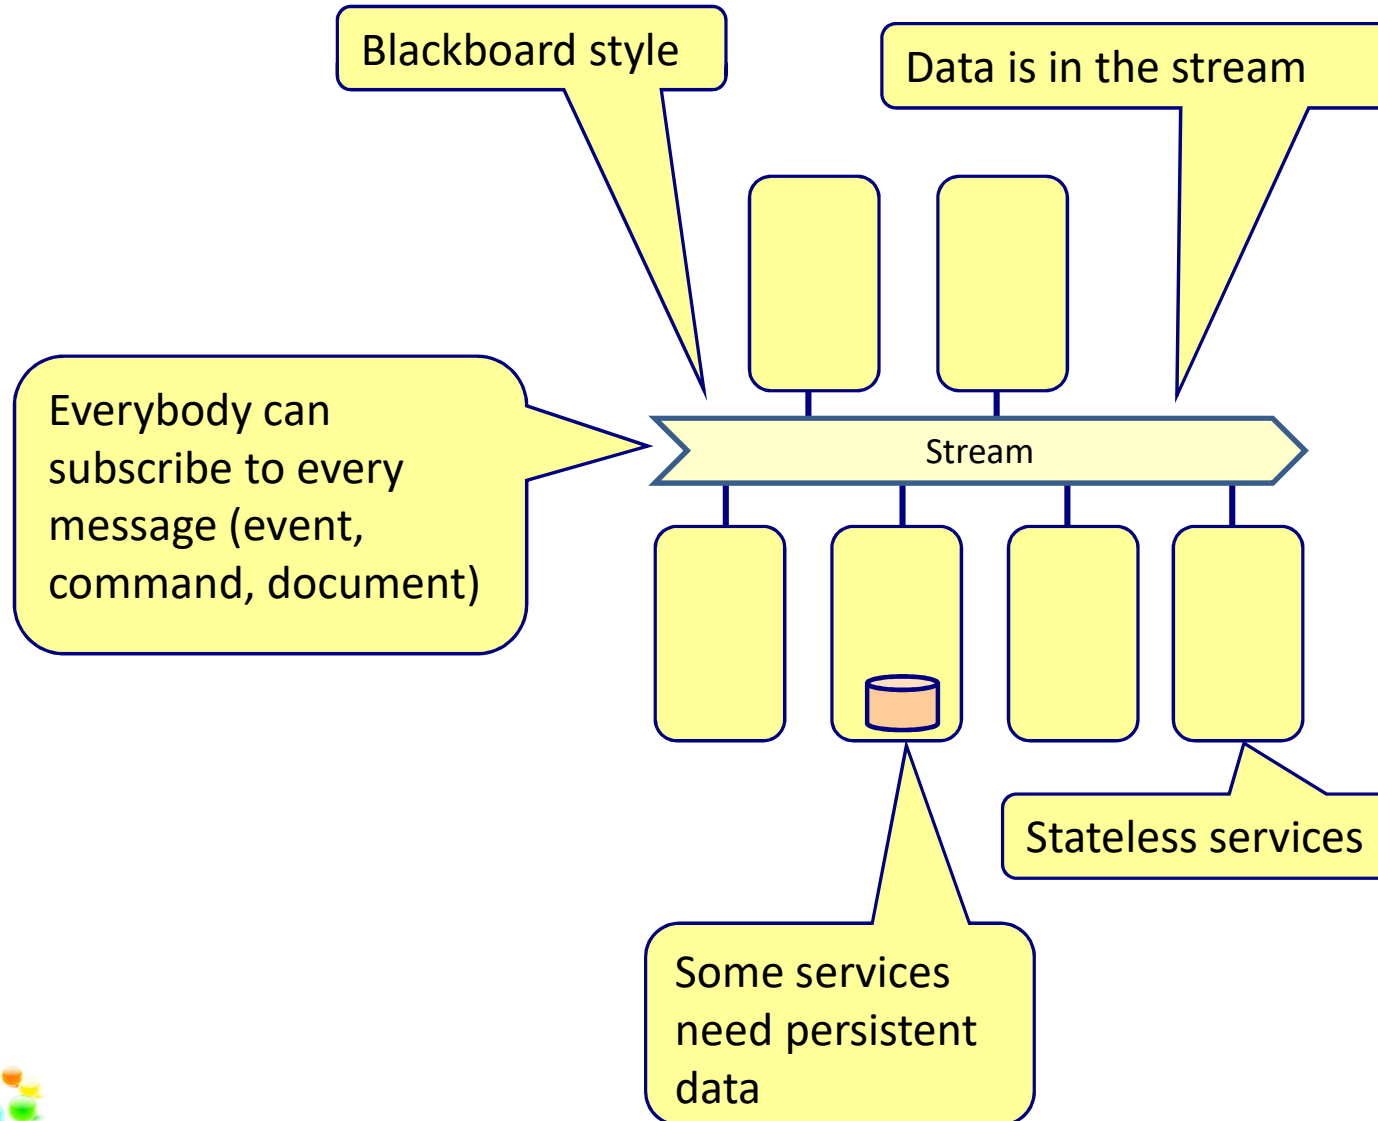


Stream processing


- Handle the data when it arrives
- Handle event (small data) by event
- Latencies in seconds or milliseconds



Where is the data?



Publish-subscribe and event sourcing



Stream based architecture

